

The Retro Playbook

Using AI to Run Better Agile Retrospectives

Prab Mutti

retro-generator.com

The Retro Playbook

Using AI to Run Better Agile Retrospectives

Copyright © 2026 Prab Mutti. All rights reserved.

No part of this book may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the author, except in the case of brief quotations embodied in reviews and certain other non-commercial uses permitted by copyright law.

This book is intended to provide general guidance on facilitating team retrospectives. The author makes no warranties as to the completeness or accuracy of its contents. The named research studies and frameworks referenced throughout are credited in the References; all example teams, names, and dialogue are illustrative.

The AI prompts in Part IV are provided as starting points. Always review AI output before sharing it with your team, and follow the privacy ground rules in Chapter 15.

Companion tools and generators: **retro-generator.com**

First edition, 2026.

Start Here

This is a field guide, not a textbook. It's built to be flipped through five minutes before a meeting as much as read cover to cover.

If you have a retro on your calendar this week, here's the fastest way in:

- **Never run one before?** Chapter 2 has a complete, minute-by-minute script. Steal it verbatim.
- **Retro's gone stale?** Jump to the format cookbook (Part III) and the icebreaker box (Chapter 4), and pick something new for what your team actually needs this week – Chapter 5 helps you choose.
- **Facilitating a tricky room?** Chapter 3 is the craft nobody teaches you: the over-talker, the silent one, the fight that breaks out, the pause you have to learn to sit in.
- **Here for the AI angle?** Part IV is written to make sense cold. One rule holds the whole thing together: *AI handles the paperwork, humans handle the talking.*

Every chapter ends with a one-page **Facilitator's Cheat Sheet** – the whole thing at a glance. Watch for recurring boxes along the way: **From the Facilitator's Chair** (what to actually say), **Try This at Home** (the same technique, applied to your own week), and **AI Shortcut** (a copy-paste prompt that does the prep for you).

You don't need to be a naturally charismatic facilitator. You need a room designed so honesty is the easy option – and that's a craft, learnable one retro at a time.

Go run one.

Contents

Part I · Retro Basics (Without the Boring Bits)

- 1 What a Retro Actually Is (and the 5 Ways They Go Wrong)
- 2 The Recipe: 5 Stages of Every Good Retro
-  *The One-Question Warm-Up*
- 3 Facilitation 101: The Stuff Nobody Tells You

Part II · The Icebreaker Box

- 4 The Icebreaker Box
-  *Rapid-Fire Facilitation*

Part III · The Retro Format Cookbook

- 5 Pick Your Format: A Decision Guide
-  *The Right Retro for the Job*
- 6 The Classics, Done Well
- 7 Formats for Specific Moments
- 8 Creative & High-Energy Formats
- 9 Remote & Async Formats
- 10 From Talk to Change: Action Items That Actually Happen

Part IV · Your AI Co-Pilot

-  *The Retro That Almost Wasn't*
- 11 Where AI Helps, Where It Hurts
- 12 AI Before the Retro (Prep in 5 Minutes)
- 13 EARS: Choosing Themes People Actually Care About
-  *The Retro That's Yours to Run*
- 14 AI During and After
- 15 Keep the Trust: Privacy Ground Rules

16 Retros When AI Is on the Team
17 Closing: The Reflective Habit

Back Matter

– Your Virtual Toolkit: Miro, Mural, and Shape Libraries
– Beyond the Retro: Building Your Own Scrum Agents
– References

PART I

Retro Basics (Without the Boring Bits)

*gets you running a decent retro by
chapter 3*

What a Retro Actually Is (and the 5 Ways They Go Wrong)

The 10-minute version of retro theory, told through five characters you already know.

The short version: a retrospective is the one recurring meeting where a team stops shipping for 45 minutes and asks "is the way we're working actually working?" That's it. No sprint is so packed that this question doesn't deserve an answer.

The long version is the rest of this book. But before any of the formats or facilitation tricks make sense, you need to see the five ways retros quietly go wrong – because odds are your team is running one of them right now, and nobody's said so out loud.

Picture five teams. Different companies, different codebases, same Tuesday afternoon slot on the calendar. Each one is about to have a retro. Here's how it goes.

The Cast of Characters

five failure modes, and the fix for each

01

The Complaint Circle

Everyone gets a turn to vent. Nobody proposes anything. The meeting ends when the clock does, not when the room feels finished.

WHY IT HAPPENS

Venting feels productive in the moment – it's real relief – so the facilitator lets it run rather than steering toward "okay, so what do we do about it?" Over time the team learns retros are where you complain, not where things change.

THE FIX

For every issue raised, ask "what's one thing that would make this 10% better?" before moving on. You don't need a full solution – you need the room to practice thinking past the complaint.

02

The Silent Room

The facilitator asks "so, how was the sprint?" and is met with the sound of Slack notifications. Two people answer every question. Everyone else nods.

WHY IT HAPPENS

Silence usually isn't apathy – it's a room where speaking up has felt costly before, or where there's no structure forcing everyone to contribute. Open questions to a group default to the two most confident voices.

THE FIX

Stop asking open questions to the whole room. Use simultaneous written input (stickies, a shared doc, silent brainstorming) so everyone answers

before anyone talks – see the icebreakers and formats chapters for exactly how.

03

The Groundhog Day Retro

"We need better documentation." "Standup runs long." "QA finds bugs too late." Said again. Same words as three sprints ago. Nobody in the room seems to notice.

WHY IT HAPPENS

Nobody's keeping a memory of past retros, so every session starts from zero. Without a written record, patterns are invisible – it just feels like déjà vu nobody can quite place.

THE FIX

Start every retro by glancing at the last one's action items for 90 seconds. "Did we do the thing we said we'd do?" This single habit kills most Groundhog Day retros – see Chapter 10.

04

The Action-Item Graveyard

The retro produces a beautiful list of six improvements. All six get written down. None get owners. None get done. The list quietly dies in a doc nobody reopens.

WHY IT HAPPENS

Six good ideas feels like a productive retro, but six items with no owner is six items nobody's accountable for. Ambition without ownership is just a wish list.

THE FIX

Leave with exactly ONE action item, one named owner, and a plan to check it at the next retro. One thing done beats six things discussed – Chapter 10 is the whole playbook for this.

05

The Blame Session

Someone says "well, if QA had actually tested this – " and the temperature in the room drops ten degrees. People start choosing words carefully. Nobody mentions their own mistakes for the rest of the meeting.

WHY IT HAPPENS

One un-checked personal jab teaches everyone in the room that honesty here is risky. After that, people protect themselves instead of the team – which is the opposite of what a retro is for.

THE FIX

Read the Prime Directive out loud, every single time, even when it feels cheesy (Chapter 2 explains why). And when blame surfaces, redirect immediately: "let's talk about the process gap, not the person."

Every one of these is fixable with structure, not personality. You don't need a naturally charismatic facilitator - you need a room designed so honesty is the easy option.

Why this book exists

Notice something about all five failure modes: none of them are about a lack of retro *formats*. Teams stuck in the Complaint Circle don't need a new activity – they need a facilitator who redirects toward action. Teams in the Silent Room don't need more enthusiasm – they need simultaneous input instead of open-floor questions.

That's the real job of this book. Chapters 2 and 3 give you the structural fundamentals – the five-stage recipe and the facilitation moves that make honesty the path of least resistance. Chapter 4 (which you may have already read) and the format cookbook in Part III give you the raw material: icebreakers and formats matched to whatever your team needs that week. Part IV shows you where AI can quietly take the admin off your plate – mining tickets for talking points, remembering what past retros raised – without ever getting near the conversation itself, because the conversation is the one part of this a machine should never run.

TRY THIS AT HOME

The Friday 15

Block 15 minutes every Friday afternoon – shorter than your commute, less effort than deciding what's for dinner. Ask yourself three questions: What went well this week? What didn't? What's the one thing I'll do differently next week?

Write the answers down somewhere you'll actually see again – a note, a journal, a running doc. The value isn't in any single Friday; it's in reading last month's answers and noticing the pattern you couldn't see in the moment. It's a personal Groundhog Day check, and it takes less time than doomscrolling.

FACILITATOR'S CHEAT SHEET

The five failure modes, at a glance

- **Complaint Circle:** venting with no push toward action. Fix – always ask "what's one thing that helps?"
- **Silent Room:** two voices dominate, the rest nod. Fix – simultaneous written input, not open questions.
- **Groundhog Day Retro:** the same issues, sprint after sprint. Fix – review last retro's action items first.
- **Action-Item Graveyard:** good ideas, no owners, nothing changes. Fix – one item, one owner, one check-in.

- **Blame Session:** honesty dies the moment someone gets personally attacked. Fix – the Prime Directive, read aloud, every time.
- **The underlying truth:** none of these are solved by a better icebreaker. They're solved by structure that makes honesty the easy option – which is what Chapters 2 and 3 build.

The Recipe: 5 Stages of Every Good Retro

Why the structure matters, how loose you can play it, and the exact script for your first one.

The recipe: Set the Stage → Gather Data → Generate Insights → Decide What to Do → Close. Every good retro format – from Mad/Sad/Glad to the fanciest remote whiteboard exercise – is a costume worn over this same five-stage skeleton, first laid out by Esther Derby and Diana Larsen [Derby & Larsen, 2006] and now the standard structure across the retrospective facilitation world.

Skip a stage and something specific breaks. Learn what each one is actually for, and you can improvise confidently instead of just following a template.

The five stages

Five words, if you're skimming: Stage, Gather, Insight, Decide, Close. Here's what each one actually buys you, and what breaks the moment you skip it.

1. Set the Stage	Turns a group of tired people into a room ready to be honest. This is your icebreaker's job – see Chapter 4. Skip it and you're asking for vulnerability from people who haven't warmed up.
2. Gather Data	Get the facts and feelings about the sprint onto the table – stickies, a board, a shared doc. Everyone contributes before anyone debates [Diehl & Stroebe, 1987]. Skip straight to discussion and the loudest opinion wins by default.
3. Generate Insights	Ask "why" of the data you just gathered. Cluster it, look for patterns, find the root cause behind the symptom. Skip this and you're solving surface complaints instead of underlying problems.
4. Decide What to Do	Narrow everything down to one committed action with one owner. This is the stage most retros skip or rush – and skipping it is exactly how you get the Action-Item Graveyard from Chapter 1.
5. Close	End on purpose – thank the room, restate the one commitment, appreciate someone. Chapter 4's Closing Check-Outs deck has quick ways to do this in under two minutes. Skip it and the retro just trails off, and trailing off is what makes people dread the next one.

You can spend 90 seconds on a stage or 15 minutes on it. You can rename it, theme it, dress it up as a sailboat or a weather map. What you can't do is skip one and expect the same result.

The Prime Directive

Before anything else, read this out loud – or say it in your own words – at the start of every retro, especially the first one:

SAY THIS, EVERY TIME

"Regardless of what we discover, we understand and truly believe that everyone did the best job they could, given what they knew at the time, their skills and abilities, the resources available, and the situation at hand."

That's the original wording from Norm Kerth [Kerth, 2001], who coined it decades ago, and it still works because it does one specific job: it tells the room, out loud, that this meeting is not a trial. It gives people explicit permission to say "I got that wrong" without it becoming an admission used against them later. Skip it and you risk the Blame Session from Chapter 1 – one unchecked jab and the room quietly decides honesty here is a liability.

Yes, it can feel a little cheesy the first few times. You will say it to a room that already trusts each other completely, and it will feel unnecessary. Say it anyway – the cheap insurance is worth it for the one retro where it isn't unnecessary. The concept it's building has a name – psychological safety, the

shared belief that the team is safe for interpersonal risk-taking – and it's not soft or fuzzy as far as the research goes: it was the single strongest factor Google found across 180 teams in its multi-year study of what makes teams effective [Google re:Work, 2015], built on Amy Edmondson's original research on the concept [Edmondson, 1999]. Stronger than who's actually on the team. One repeated, slightly-cheesy sentence is a remarkably cheap way to buy a share of that.

Your first retro, minute by minute

If you've never facilitated one before, here's the whole thing scripted. Steal it verbatim for retro number one; improvise more from retro number two onward.

0:00 – OPEN

"Thanks everyone for being here. Quick reminder before we start – [read the Prime Directive]. "This isn't about blame, it's about what we learn and improve. Let's begin with a quick check-in."

0:03 – SET THE STAGE

Run a 2-minute icebreaker (Chapter 4 has 29 of them – Fist to Five is a safe first pick). "Alright, that gives us a read on the room. Let's get into it."

0:08 – GATHER DATA

"For the next 8 minutes, silently write down what went well and what didn't – one item per sticky or line. Don't discuss yet, just write." Set a visible timer. Silence is fine; let it sit.

0:16 – GENERATE INSIGHTS

"Let's group similar items together." Cluster the input as a group. Then: "Looking at these clusters, what's the pattern underneath? Why do we think this keeps happening?"

0:30 – DECIDE

"Of everything we've discussed, what's the ONE thing we commit to changing before next retro? Who can own that?" Push for a single, specific, owned action – resist the urge to write down all six good ideas as commitments.

0:40 – CLOSE

"So we're committing to [X], owned by [name], and we'll check it next time. Before we go – does anyone want to appreciate someone for something this sprint?" Let one or two people answer. "Thanks everyone, see you next sprint."

Forty-five minutes, five stages, one commitment. That's a complete retro – everything else in this book is variation on this theme.

TRY THIS AT HOME

The five stages, for anything

This structure isn't retro-specific – it's just how humans reflect well. Planning a family holiday debrief? Set the stage (pour a drink), gather data (what actually happened), generate insights (why did the airport day go sideways), decide (book earlier flights next time), close (say what you're grateful for about the trip). Same skeleton, completely different room.

FACILITATOR'S CHEAT SHEET

The recipe, in one box

- **Set the Stage:** warms the room up before asking for honesty. Never skip, even if it's 90 seconds.
- **Gather Data:** everyone contributes before anyone debates. Written and simultaneous beats spoken and sequential.
- **Generate Insights:** ask "why" of what you gathered – symptoms vs. root causes.
- **Decide:** ONE action, ONE owner. This is the stage most retros rush – don't.
- **Close:** end on purpose. A trailed-off retro is why people dread the next one.

— **Say the Prime Directive out loud, every time.**

It's the sentence that makes stage 2 possible.

The One-Question Warm-Up

No quiz, no answer key - just one job. Pick a single question, open the room with it, and get every voice making a sound before the real work starts.

How to use it: choose one line that fits the mood, ask it, and go round the room (or drop answers simultaneously in chat). One word or one sentence each - no speeches. The answers are quietly data: three storm-cloud check-ins in a row and you already know today's real topic.

Read the temperature

- 1 Describe this sprint in one word - no explanation unless a word begs for one.
- 2 What's your internal weather right now: sunny, cloudy, foggy, thunderstorms?
- 3 Fist to five - how was the sprint? Show your hand on three.
- 4 What percent is your battery at, and what would charge it?

Lift a flat room

5 If this sprint were a movie title, what would it be?

6 What superpower did you bring this sprint, and what was your kryptonite?

7 Post the one GIF that best sums up your fortnight.

Build a little trust

8 Appreciate one teammate for something specific they did this sprint.

9 What was your very first job, and what did it teach you?

10 One thing you messed up this sprint - then take a bow while we applaud.

Pick one, mean it, and always allow a pass. An icebreaker someone survived is worse than one they skipped.

Want the full set - 29 icebreakers sorted by exactly this mood-first logic, plus closing check-outs for the end of the room? That's Chapter 4, the Icebreaker Box.

Facilitation 101: The Stuff Nobody Tells You

The real craft – what to actually say and do in the ten moments that make or break a retro.

Nobody hands you a script for this. You get told to "facilitate the retro" and left to figure out, live, in front of your team, what to do when Dave won't stop talking or the room goes dead silent. This chapter is that script – ten recurring moments, and exactly what to say in each one.

Ten Moments You'll Face

what to actually do, and the line to have ready

01

Opening the Room

The first thirty seconds set the tone for the whole meeting. Walk in flustered and rushed, and the room mirrors it back at half your energy.

WHAT TO DO

Arrive on time, greet people by name as they join, and open with calm, deliberate pacing – even if you're not feeling calm. The room takes its emotional cue from you before a single word of content happens.

TRY THIS LINE

"Hey everyone – before we dive in, take a breath, this is a no-judgment space" – then go straight into the Prime Directive from Chapter 2.

The Over-Talker

One person answers every question first, at length, and by the third round everyone else has mentally checked out and started deferring to them.

WHAT TO DO

Don't shut them down – redirect the room's attention structurally. Switch to written/simultaneous input (Chapter 2's Gather Data stage) so their volume can't dominate the format itself. If it's a live discussion, name a round: "let's hear from folks who haven't spoken yet."

TRY THIS LINE

"Great point – let's park that and hear a couple more perspectives before we go deeper." Said warmly, not as a rebuke.

The Silent One

Someone hasn't said a word in twenty minutes. You don't know if that's comfort, disengagement, or something worse – and calling them out cold feels like a spotlight, not an invitation.

WHAT TO DO

Never put someone on the spot in front of the group. Give a low-stakes, specific opening instead – pointing at something they wrote, not asking them to generate content live. Silence in a retro isn't always a problem to solve mid-meeting; sometimes it's a signal to check in 1:1 afterward.

TRY THIS LINE

"[Name], I saw you wrote something about the deploy process – want to say more about that?" Specific and low-pressure beats "what do you think?"

04

Sitting With Silence

You ask a question. Nobody answers immediately. Four seconds feels like an hour, and every instinct screams to fill the gap yourself.

WHAT TO DO

Count to ten in your head before you say anything else. Genuinely – this is the same finding as classroom "wait time" research: teachers typically wait under a second for a student to answer, and simply extending that pause to three-plus seconds measurably improves the length and quality of the response [Rowe, 1972]. A retro room works the same way. The pause is doing its job; don't rescue the room from it.

TRY THIS LINE

Nothing. That's the whole skill – say nothing, hold eye contact around the room, and let it breathe.

05

Staying Neutral

You have a strong opinion about the thing being discussed – you were IN the incident, you know exactly whose fault the outage was – and it's dangerously easy to steer the retro toward your version of events.

WHAT TO DO

The facilitator's job is to run the process, not win the argument. If you must contribute content, say so explicitly and step out of the facilitator role for a moment, then step back in. Ask questions more than you make statements.

TRY THIS LINE

"Stepping out of facilitator mode for a second – my take is X. Okay, back to the room: what does everyone else think?"

Timeboxing Without Being a Robot

The timer hits zero mid-sentence during the most honest moment of the retro. Cutting it off kills the value; ignoring the clock kills every retro after this one when it runs 90 minutes.

WHAT TO DO

Set visible timers so the room shares the constraint with you – it stops feeling like your arbitrary rule. When something genuinely important is unfolding, extend consciously and say so out loud, then cut the next, less critical stage instead.

TRY THIS LINE

"This is important, let's give it two more minutes – and we'll trim the wrap-up to make up the time."

AI SHORTCUT

If your org has it, this is one job AI is genuinely good at handing off entirely: Microsoft Teams' Facilitator agent will run a visible countdown per agenda topic, alert the room when time's up, and nudge everyone at the halfway point and before wrap-up [[Microsoft, Teams Facilitator docs](#)] – built-in for remote retros, no separate timer app needed.

07

Abandoning Your Plan

You planned a Sailboat retro. Ten minutes in, it's obvious the room actually needs to talk about the incident from Tuesday, and the format you picked is actively getting in the way.

WHAT TO DO

The format serves the room, not the other way around. A good facilitator can feel the mismatch and pivot – drop the metaphor, ask the real question directly, come back to structure once the real conversation has had air.

TRY THIS LINE

"Let's set the format aside for a minute – it sounds like we need to talk about Tuesday. What actually happened?"

08

Whole-Room Disengagement

This isn't one silent person – it's everyone. One-word answers, phones face-up on the table, the same three people going through the motions. The room isn't hiding something; it's just not here.

WHAT TO DO

Don't push harder into the agenda you planned – that's fighting the symptom, not finding the cause. Name it and ask the room, anonymously, where they actually are: Derby and Larsen's ESVP check has people silently mark themselves as an Explorer (here to dig in), Shopper (here for one useful idea), Vacationer (happy for the break, not engaged) or Prisoner (would rather be anywhere else) [Derby & Larsen, 2006]. A room full of Vacationers and Prisoners isn't a facilitation failure to fix in the next ten minutes – it's information. Shorten what's left, or ask directly why the energy is where it is.

TRY THIS LINE

"I'm sensing low energy in the room today – quick anonymous check: on a scale of Explorer to Prisoner, where are you at with this retro right now?"

09

When It Turns Into a Fight

Two people are relitigating something that happened weeks ago. Voices are rising, the jabs are getting personal, and the rest of the room has gone quiet in the specific way that means they're watching, not participating.

WHAT TO DO

Don't pretend it isn't happening and don't pick a side – name it out loud, calmly, and pause the exchange rather than letting it run. Lencioni's distinction is the useful one here: conflict about an idea or a decision is healthy and worth having; conflict that's become about the person isn't, and the facilitator's job is to pull it back toward the first kind [Lencioni, 2002]. If it keeps sliding back into personal territory even after you name it, don't try to resolve it live in front of the whole team – park it and take it to a 1:1.

TRY THIS LINE

"Let's pause there – I can tell this matters to both of you. Let's separate what we actually disagree about from the history, and if it needs more time than we have, I'll set up time for the two of you after this."

The Conversation That Won't Stay on Track

Someone raises a real point that isn't this retro's topic – a tooling gripe, a roadmap question, an argument from three sprints ago – and the room follows it down the tangent. Or worse: the current topic itself has gone in a circle twice already and nobody wants to be the one to call it.

WHAT TO DO

Two different tools for two different problems. For the tangent that's genuinely off-topic but worth remembering, keep a visible

CAR PARK

(a column on the board, a sticky note area) and move it there in front of everyone – it's not being dismissed, it's being deferred, and writing it down is what makes that credible. For a conversation that's *on*-topic but circling without landing anywhere, agree the ground rule up front that anyone – not just the facilitator – can call

ELMO

("Enough, Let's Move On"), a one-word interrupt the whole team is allowed to use without it being a slight on whoever was talking.

TRY THIS LINE

"That's worth talking about – let's Car Park it and come back after the retro." Or, if you've agreed the ground rule in advance: "I'm going to call ELMO on this one – we've circled it twice, let's move to the next thing."

None of this is about charisma. It's ten specific, learnable moves - and like any craft, they get easier the fifth time than the first.

Reading the room's energy

Before you facilitate a single stage, spend the first sixty seconds just watching. Are people leaning in or checking phones? Is the energy flat, tense, or oddly upbeat? This single read should decide which icebreaker you reach for (Chapter 4 is organized exactly by this – low energy needs a lift, guarded rooms need a trust builder) and whether your planned format still fits. Facilitation isn't executing a script perfectly; it's executing a script while constantly checking whether it's still the right one.

Why the facilitator shouldn't always be the Scrum Master

The default on most teams is that the Scrum Master runs every retro, every time, forever. It's convenient – they already own the calendar invite – but it quietly costs the team more than it saves. Three specific reasons to rotate the role instead:

- **The Scrum Master often has a stake in the answer.** If the retro is about whether the process the Scrum Master owns is working, having them also run the meeting that judges it is an obvious conflict – not of bad intent, just of position. A rotating facilitator has less riding on any particular conclusion.

- **Facilitation is a skill worth spreading, not hoarding.** Every skill in this chapter – reading a room, sitting with silence, redirecting an over-talker – is useful far beyond retros. A team where only one person ever practices it is a team where that skill lives in one head.
- **A fresh facilitator asks different questions.** The person who ran the last six retros has patterns and blind spots baked in by now. Someone facilitating for the first time in months asks the "obvious" question a practiced facilitator has stopped noticing is worth asking.

In practice: rotate on a fixed cadence (every retro, or every few) rather than leaving it to whoever volunteers – volunteering skews toward the people already comfortable with the role, which defeats the point. Give the outgoing facilitator five minutes beforehand to hand off anything the next person needs to know, and don't expect the first attempt to be as smooth as the Scrum Master's – that's the cost of building the skill across the team, and it's worth paying.

Keep an ear to the ground

Sir Alex Ferguson managed Manchester United for 26 years and won 38 trophies doing it, and one habit shows up again and again in accounts of how he actually ran things: he was relentless about gathering information from everyone around him, all the time, not just in the moments people expected him to be paying attention. After matches, he'd go into the opposing manager's office specifically to listen to the rival coaching staff – reportedly picking up on players like Andy Cole, Michael Carrick, and Rio Ferdinand that way. He asked his own players who impressed them on the pitch, which is part of how United ended up signing Cantona, Keane, Rooney, and Ronaldo. Harvard Business School ran a full case study on his methods after watching him work for a season

[Elberse & Ferguson, 2013], and in his own memoir he put the underlying philosophy plainly:

"There's a reason that God gave us two ears, two eyes and one mouth. It's so you can listen and watch twice as much as you talk."

A facilitator's version of this is smaller in scale but the same idea. The retro isn't the only place you learn what a team actually cares about – standups, water-cooler chat, an offhand comment in a Slack thread are where you pick up the nugget that a good icebreaker or a themed retro (Chapter 12's "Theme It" prompt) can turn into something the whole team enjoys. Nobody announces "by the way, half of us have been obsessed with the World Cup this month" in a formal meeting. You find that out by listening the other 29 days of the sprint, not the one day you're facilitating. As Ferguson put it in the same book, "a network takes time to develop" – built from repetition, and from how you treat people in between the moments that matter [Elberse & Ferguson, 2013].

AI can approximate a slice of this – mining Jira comments or Slack threads for what a team's been talking about, the same way Chapter 12's prep prompts mine tickets for what happened. It's worth trying. But it can't replace being in the room for the comment that never gets typed anywhere – the sigh after a stand-up, the joke that lands a little too well, the thing someone only says once, to you, in passing. That's the human touch from Chapter 11's rule, applied to information-gathering instead of facilitation itself: AI can help you notice patterns in what's written down. It can't be the ear to the ground.



TRY THIS AT HOME

The pause, off the clock

The "sit with silence" skill is worth stealing for any hard conversation – with a partner, a friend, your kid. Ask the real question, then physically count to ten before you speak again. Most people's honest answers arrive a beat after the socially comfortable one does.

FACILITATOR'S CHEAT SHEET

Ten moments, ten moves

- **Open calm** – the room mirrors your energy in the first thirty seconds.
- **Over-talker:** switch to written/simultaneous input rather than calling them out.
- **Silent one:** a specific, low-pressure invite beats an open "what do you think?"
- **Silence:** count to ten before you rescue the room from it.
- **Neutrality:** ask more than you state; step out of "facilitator" explicitly if you must weigh in.
- **Timebox:** visible timers, and extend consciously out loud when it's worth it.
- **Abandon the plan** when the room needs something the format is blocking – the format serves the room, not vice versa.

- **Whole-room disengagement:** an anonymous ESVP check names low energy without putting anyone on the spot.
- **A fight breaks out:** name it, don't pick a side, and pull it back from personal to the actual disagreement – or take it offline.
- **Off-track conversation:** Car Park a genuine tangent so it's deferred, not dismissed; agree ELMO as a team ground rule for cutting a circular one short.
- **Rotate the facilitator.** Don't leave it permanently with the Scrum Master – it's a conflict of interest when the retro judges their own process, and the skill is worth spreading across the team.
- **Keep an ear to the ground.** The best material for your next retro or icebreaker gets picked up in standups and hallway chat, not invented cold five minutes before the meeting.

PART II

The Icebreaker Box

one page per icebreaker, flip and go

The Icebreaker Box

29 icebreakers and check-outs that don't make people cringe – plus a would-you-rather bank for reviving a room mid-retro – sorted by what your team needs right now, not the alphabet.

Why bother: a good icebreaker gets every voice in the room to make a sound in the first five minutes. It's a facilitator's rule of thumb, not a lab result, but it tracks a real and well-documented pattern: groups tend to suppress minority or early-uncertain views once a few confident voices set the tone[Noelle-Neumann, 1984] – so getting everyone talking before that dynamic sets in matters. That's the whole job of an icebreaker – it's the ignition, not the fluffy bit before the real meeting.

How to use this chapter: don't read it start to finish. Flip to the section that matches your team's mood, pick one card, done. Rotate – even the best icebreaker dies after four consecutive uses.

A note on time: every format and icebreaker in this book, unless a card says otherwise, is built to fit inside a normal 45-minute-to-1-hour retro – that's the length most teams actually run, and it's the length this book is designed around. The one exception is the Full Activities section at the end of this chapter: longer, classroom-style exercises that teach a concept rather than gather sprint data, and that need a session booked specifically for them.

||| Energy – pips filled =
livelier

• Remote-friendly
/

•• built for
remote

🕒 Time · Group size

Quick Temperature Checks

under 2 minutes – for tight schedules or teams who find icebreakers cheesy

One-Word Check-In

01

🕒 1–2 min any ||| ⚙️

Go around the room: "Describe this sprint in one word." No explanations, no follow-ups – unless someone's word begs for one ("Interesting, Dave? We'll come back to that").

WATCH OUT

Enforce ONE word. The moment someone gives a paragraph, everyone after them will too.

Weather Report

02

🕒 1–2 min any ||| ⚙️

"What's your internal weather right now?" Sunny, cloudy, foggy, thunderstorms. The metaphor gives cover to be honest – "a bit foggy" is easier to say than "I'm burnt out and confused."

WATCH OUT

If someone says "thunderstorm," don't interrogate them publicly. Note it, check in privately, or let the retro surface it naturally.

Emoji Check-In

03

🕒 1 min any ||| 🎲 🎲

Everyone posts one emoji describing their sprint in chat, all at once (3, 2, 1, post). Then eyeball the spread: mostly ? There's your retro agenda.

WATCH OUT

Simultaneous reveal matters. Posted one at a time, everyone anchors on the first emoji.

Fist to Five

04

🕒 1 min any ||| 🖐️

"How was the sprint, fist to five?" Everyone shows fingers simultaneously – fist is terrible, five is glorious. Instant histogram of the room. Ask one low scorer and one high scorer for a sentence each.

WATCH OUT

This is a conversation starter, not a metric. The moment someone puts the average in a spreadsheet, honesty dies.

Battery Check

05

🕒 1–2 min any ||| ⋮

"What % is your battery at, and what would charge it?" The second half is the payload – "20%, and a weekend without pager duty" tells you more than any survey.

WATCH OUT

If the whole team is under 40%, stop. That IS the retro topic now.

Energy Lifters

3–8 minutes – for the flat room, the sprint-from-hell aftermath, or 4pm on a Friday

My Sprint as a Movie Title

06

🕒 3–5 min 3–12 ||| ⋮

Everyone invents a movie title for the sprint. "Gone in 60 Deployments." "The Fast and the Spurious." Laughter guaranteed, and the titles are secretly data – three disaster movies in a row is a signal.

WATCH OUT

Give people 60 seconds of silent thinking first. Put on the spot, half the room blanks.

Two Truths and a Lie: Work Edition

07

🕒 5–8 min 3–8 ||| ⋮

Each person states three things about their sprint; the team guesses the lie. Works because the truths are often more surprising than the lie.

WATCH OUT

Caps at ~8 people or it eats your retro. For bigger groups, ask for two volunteers.

Desert Island Tools

08

🕒 3–5 min any ||| ⋮

"You're stranded on a desert island and can keep only one tool from our stack. What survives, and what are you secretly glad drowned?" The second half is where the honesty lives.

WATCH OUT

Don't let it become a serious tooling debate. Park real grievances on the board for the gather-data stage.

GIF Battle

09

🕒 3–4 min any ||| ⚫ ⚫

"Post the GIF that best describes this sprint." Optional: vote for a winner, who gets nothing but glory. Fastest route from silence to laughter in a remote call.

WATCH OUT

Set a 2-minute timer or people will hunt for the perfect GIF for ten.

Superpower / Kryptonite

10

🕒 4–6 min 3–10 ||| ⚫

Everyone names the superpower they brought this sprint and the kryptonite that weakened them. Fun wrapper, real data – the kryptonites often become the retro's main course.

WATCH OUT

People underclaim their superpower. If someone says "nothing really," let a teammate answer for them.

Sprint Soundtrack

🕒 3–5 min any ||| ⋮

"If this sprint had a theme song, what would it be?" From "Eye of the Tiger" to the Benny Hill theme. Power move: build a team playlist over the quarter as an audio mood chart.

WATCH OUT

Nothing kills it faster than pausing to actually play each song. Titles only.

Would You Rather

a bank of forced-choice questions – as much a mid-retro rescue as an opener

Why they work: a binary choice removes the blank-page tax. There's no clever answer to compose and nothing to get wrong, so even the person who dreads open questions can pick a side in half a second – participation is instant and total. And a forced trade-off is quietly revealing: ask a team "what do you value?" and you get platitudes; ask "would you rather cut the feature or cut the deadline?" and you find out what they actually optimise for, because now they have to choose.

Not just an opener: most icebreakers earn their keep in the first five minutes and then get out of the way. Would-you-rather travels. Drop a round in the *middle* – when a thread turns heated, when the gather-data stage has flattened into a slog, or at the 40-minute mark when the room's battery hits zero. Ninety seconds of duck-sized horses is a pattern-interrupt: it lets everyone exhale, breaks the spiral, and you re-enter the real work lighter. Keep two or three in your back pocket for exactly the moment a retro starts going south.

How to run it: read the question, everyone picks a side at once – hands up, thumbs, A/B in chat, or physically step to the left or right of the room. Then ask one voice from each side a single "why?" That "why" is the whole payload; the vote is just the excuse to get there. One or two questions is plenty – timebox hard and get out while it's still fun.

Pure Nonsense

Zero work content, nothing to be right or wrong about – reach for these when the room is flat or fried and you need laughter fast.

🕒 2–3 min any III

Would you rather...

- ...fight one horse-sized duck, or a hundred duck-sized horses?
- ...have to sing everything you say, or dance everywhere you walk?
- ...be able to teleport but only to places you've already been, or fly but no faster than walking pace?
- ...have fingers as long as your legs, or legs as short as your fingers?
- ...live somewhere it's permanently 5°C, or permanently 30°C?
- ...only ever be able to whisper, or only ever be able to shout?
- ...have a dramatic slow-motion walk-in every time you enter a room, or background music only you can hear?
- ...have Wi-Fi that reaches everywhere but is always slow, or blazing-fast Wi-Fi in only one room of the house?

FACILITATION NOTE

The debate is the point, not the answer. Ask two people on opposite sides to defend themselves for ten seconds each – the more ridiculous the justification, the faster the room thaws.

The Developer's Dilemma

Same low stakes, but the jokes land closer to home – good for a warmed-up room that enjoys a bit of shared suffering.

🕒 3–4 min any 

Would you rather...

- ...deploy on a Friday afternoon, or review a 2,000-line pull request?
- ...debug someone else's regex, or someone else's CSS?
- ...inherit undocumented code that works, or documented code that doesn't?
- ...live with tests that fail 1% of the time for no reason, or no tests at all?
- ...give up your standups, or give up your Slack notifications?
- ...have every meeting on your calendar for a week, or your IDE stuck in light mode forever?
- ...merge without a code review, or wait three days for one?
- ...have perfect documentation but no Google, or Google but no documentation?
- ...chase the same tiny bug forever, or ship one huge bug you can never find?

FACILITATION NOTE

These double as a pressure valve for real gripes. If "flaky tests vs no tests" gets genuinely heated, you've just found a gather-data topic – park it on the board and move on.

Secretly Data

Trade-offs that look like a game but surface what your team actually optimises for. Reach for these when a discussion is stuck in vague agreement and you want the real priorities on the table.

🕒 3–5 min any 

Would you rather...

- ...ship fast and fix in production, or ship slow and ship it right?
- ...cut the feature, or cut the deadline?
- ...have a smaller scope done well, or the full scope done roughly?
- ...have perfect requirements with a tight deadline, or vague requirements with all the time you need?
- ...a teammate tell you your idea is bad to your face, or say nothing and quietly work around it?
- ...over-communicate and risk noise, or under-communicate and risk surprises?
- ...ship it and apologise later, or ask permission and maybe miss the window?
- ...have a teammate who's brilliant but unreliable, or reliable but average?
- ...fix the root cause slowly, or the symptom now and the root cause "later"?

FACILITATION NOTE

Don't rush to the "why." Let people physically commit to a side first, then read the split – if it's roughly 50/50, that disagreement *is* your retro.

The Deeper Cut

Gentle, human trade-offs for a team that already trusts each other. Best late in a good retro or as a check-out – not a cold open with strangers.

🕒 4–6 min 3–10 ■■■

Would you rather...

- ...be known as the most reliable person on the team, or the most creative?
- ...get public recognition for a team win, or a quiet thank-you for a personal one?
- ...have a job you're great at but bored by, or one you love but struggle with?
- ...always know exactly how you're doing, or be left alone to just do the work?
- ...have more time, or more autonomy?
- ...be respected by your team but not your manager, or by your manager but not your team?
- ...master a skill you'll rarely use, or be merely passable at one you use every day?

FACILITATION NOTE

Always allow a pass, and you go first. These edge toward the personal, so a short honest answer from the facilitator gives everyone permission to keep theirs light if they want to.

Office & Ritual

The everyday absurdities of how teams actually work – meetings, PTO, the office. Light, universally relatable, and a good reset when process talk turns grumpy.

🕒 3–4 min any 

Would you rather...

- ...sit through a meeting that could've been an email, or send an email that spawns a meeting?
- ...have unlimited PTO you feel guilty taking, or a fixed 20 days you take without a second thought?
- ...work four ten-hour days, or five six-hour days?
- ...never have another status meeting, or never have another surprise deadline?
- ...permanently work from home, or permanently work from the office?

FACILITATION NOTE

The splits here are quietly a culture survey. A room that unanimously picks "never another status meeting" has just handed you a retro topic.

Sci-Fi & Superpowers

Pure imagination, zero work content – the widest-open, most playful batch, for when the room needs to stop thinking about the sprint entirely.

🕒 2–3 min any III

Would you rather...

- ...pause time for everyone but yourself, or rewind ten seconds once a day?
- ...read minds but never be able to switch it off, or be invisible but only when nobody's looking?
- ...speak every human language, or be able to talk to one animal of your choice?
- ...never need sleep, or never need to eat?
- ...have a theme song that plays when you walk in, or subtitles showing what you're really thinking?

FACILITATION NOTE

No debrief needed – this one's purely to reset the mood. Short on time? Pick one, take three answers, move on.

Trust Builders

5–10 minutes – new teams, merged teams, or rooms being politely dishonest

Appreciation Shower

12

🕒 5–8 min 3–10 ||| ⋮

Each person appreciates one teammate for something specific from this sprint. Specific beats generic – "thanks for staying on the deploy call till midnight" lands; "thanks for being great" doesn't.

WATCH OUT

Make sure nobody's left out. Quietly track who's been appreciated; the facilitator fills any gap.

One Thing You Don't Know About Me

13

🕒 5–10 min 3–8 ||| ⋮

Everyone shares one non-work fact the team doesn't know. Teams that know each other as humans give each other benefit of the doubt when code review gets tense.

WATCH OUT

Always allow a pass. This must never feel like forced disclosure – that builds the opposite of trust.

Kudos Cards

14

🕒 5 min any ||| ⚙️

Everyone writes anonymous thank-you notes to teammates on stickies; facilitator reads them aloud. The anonymity lets people say warmer things than face to face.

WATCH OUT

Seed the deck. Write two or three yourself before the retro so nobody stares at an empty board.

The Failure Bow

15

🕒 5–8 min 3–10 ||| ⚙️

Borrowed from improv: each person shares one thing they messed up, then takes a theatrical bow while everyone applauds. Physically rehearses the idea that admitting mistakes is safe here.

WATCH OUT

The facilitator goes first with a real failure, not a humble-brag.

16

First Job Stories

🕒 5–8 min 3–8 ||| ⚙️

"What was your first ever job, and what did it teach you?" Instantly humanizing, zero risk, and the lessons often echo the team's current situation.

WATCH OUT

Nothing – nearly unbreakable. Just timebox it; the stories are good.

Remote-First Icebreakers

designed for distributed teams, not adapted for them

17

Camera Scavenger Hunt

🕒 3–5 min any ||| ⚙️

"You have 60 seconds: grab the nearest object that represents this sprint. Go." People return with stress balls, cold coffee, a tangled cable. Movement plus metaphor plus laughter.

WATCH OUT

Some people can't leave their desk. Offer "or find an image on your desktop" as the alternative.

Virtual Background Stories

18

🕒 3–5 min 3–10 ||| ⚫ ⚫

Before the retro: "Set a virtual background that represents your sprint."
A dumpster fire, self-explanatory. Each person explains theirs in one sentence.

WATCH OUT

Needs advance notice – put it in the calendar invite, or you'll get five default backgrounds and shrugs.

Chat-Storm

19

🕒 2–3 min any, 8+ ||| ⚫ ⚫

Ask a question, everyone types an answer, nobody presses Enter until you count down – then the chat floods at once. Simultaneity beats anchoring.

WATCH OUT

Read some answers aloud, or people learn their input vanishes into the void.

Home Desk Show-and-Tell

20

🕒 4–6 min 3–8 ||| ⬤ ⬤

One or two people per retro (rotate!) give a 60-second tour of one thing on or near their desk. Slowly builds the "we've been to each other's houses" feeling remote teams lack.

WATCH OUT

Rotation, not volunteers – the same two extroverts will otherwise host every episode.

Async Mood Board

21

🕒 0 live any ||| ⬤ ⬤

The day before, drop a board with one prompt: "Add an image, meme, or sentence for this sprint." Open the retro by silently reviewing it together for two minutes.

WATCH OUT

If the board is empty an hour before the retro, nudge the team channel with your own entry first.

Themed & Seasonal

for special occasions – used sparingly, these are the ones people remember

The Holiday Special

22

🕒 3–5 min any III 🔴

Before a break: "What are you leaving behind this year, and what are you taking with you?" Half icebreaker, half closing ceremony.

WATCH OUT

Keep it light-touch – anchor on the break, not the holiday, since not everyone celebrates the same one.

Launch Day Toast

23

🕒 3–5 min any III 🔴

Just shipped something big? Everyone raises a drink and toasts one person, one moment, or one small miracle from the release. Celebration first, surgery second.

WATCH OUT

Alcohol optional and unremarked. Coffee toasts count.

24

Post-Crunch Decompression

🕒 5 min any ||| ⚙️

After a brutal push: "What's the first non-work thing you did when it ended?" Normalizes recovery – and the person who answers "started on the next thing" needs a different conversation.

WATCH OUT

Don't let this excuse skipping a real retro about the crunch itself.

25

Sprint Horoscope

🕒 3–5 min 3–10 ||| ⚙️

Read the team a mystic "horoscope" for the sprint just gone: "Mercury was in retrograde, which explains staging. A tall dark stranger – the new API – brought chaos to your house of deployments..."

WATCH OUT

Needs 5 minutes of prep – see the AI Shortcut below, it's built for exactly this.

26

Time Capsule

🕒 5 + 5 min any ||| ⋮

Quarterly: everyone writes one prediction and one hope, sealed away. Next quarter, open last quarter's capsule first. The gap between prediction and reality is the best retro material you'll ever get free.

WATCH OUT

Calendar the opening or it never happens. The magic is entirely in the reveal.

Closing Check-Outs

under 2 minutes – for the Close stage (Chapter 2), the bookend to the temperature checks that opened the room

27

Happiness Check-Out

🕒 1 min any ||| ⋮

Everyone rates how they're leaving this retro – not the sprint, this meeting – 1 to 5 or an emoji, posted at once. It's the same move as the opening temperature checks, just pointed at the room you built instead of the sprint you're reviewing.

WATCH OUT

If the number is lower than the check-in, ask why before you close – a retro that leaves people worse off than it found them is the exact failure mode it's meant to fix.

Plus / Delta

28

🕒 1–2 min any ||| ⚙️

Two quick shouts: one thing about *this retro* that worked (Plus), one thing you'd change next time (Delta). Feedback on the format itself, not the sprint – the fastest way to know if a format's overstayed its welcome before Chapter 5 tells you to rotate it.

WATCH OUT

Keep it about the retro, not the sprint – if answers drift back to sprint content, that's the Gather Data stage leaking into Close.

Gratitude Round

29

🕒 2–3 min any ||| ⚙️

One or two people say thanks to a teammate for something specific this sprint, out loud, before anyone logs off. This is the literal script from Chapter 2's Close stage, given its own card so it doesn't get skipped when the retro runs long.

WATCH OUT

Don't force everyone to participate – two genuine thank-yous beat six half-hearted ones.

Full Activities

*book extra time for these – classroom-style exercises that teach a concept,
not quick sprint check-ins*

The Ball Point Game

Teams pass balls through a set of hand-offs as fast as they can, run several timed rounds with a short retro between each, and watch their own throughput climb – a physical, undeniable demonstration of why short iterations and a quick retrospective between them actually work.

🕒 **60–75 min** **6–12** **~100 balls, 2 baskets**

1. Explain the rules: every ball must get air-time, every team member must touch each ball, and only the person who launched a ball may land it in the basket. A ball on the floor is a defect.
2. Give the team 2 minutes to plan and estimate how many balls they'll get through.
3. Run a 2-minute timed round. Count the valid balls.
4. Give 1 minute to retro and improve the process, then run it again – five rounds total.
5. Debrief: what changed between round 1 and round 5, and why?

FACILITATION NOTE

The debrief is the whole point – don't skip it for time. Ask directly: "what did five short cycles get you that one long round wouldn't have?"

WATCH OUT

Needs open floor space and a genuine 60+ minutes – this is the classic version, credit to TeamRetro's write-up for the exact rules above.

VARIANT

Some facilitators swap the small balls for a single football, thrown or kicked hand-off to hand-off instead of tossed – same rules, same lesson, more space and louder laughter required.

The WIP Limit Game

Everyone writes three short lists – numbers 10 down to 1, letters J down to A, roman numerals X down to I – twice: once switching between all three lines after every character (multitasking), once finishing one line before starting the next (a work-in-progress limit of one). Timed both ways, the difference is dramatic and entirely self-inflicted.

🕒 25–35 min any paper, a stopwatch

1. Everyone draws three rows on paper: numbers 10 → 1, letters J → A, roman numerals X → I.
2. Round 1 (multitasking): write one character in row one, then one in row two, then one in row three, repeat until all three rows are complete. Time it.
3. Round 2 (WIP limit 1): complete row one fully, then row two, then row three. Time it.
4. Compare times and error counts out loud – round 2 is reliably faster with fewer mistakes.

FACILITATION NOTE

Let the team feel it before you explain it – reveal the "why" (context-switching cost) only after both rounds are done and times are compared.

WATCH OUT

Known in different training circles as the Context-Switching Exercise or the Personal Kanban WIP Exercise – same mechanic, several names.

Noah's Ark (Animal Sizing)

A relative-estimation exercise, not a literal game with fixed rules – the "ark" is just a fun visual frame for Animal Sizing, where work gets compared to animals of increasing size (flea, ant, bee, rabbit, dog, lion, elephant...) instead of story points, because size intuition transfers faster than a number scale does.

🕒 **30–45 min** **any** **an ark drawing, animal cards**

1. Draw a simple ark on the board (or use a template).
2. Pick 6–8 real backlog items to size.
3. For each item, the team places it into the ark next to the animal that "feels" like the same size – no discussion of hours, no debate about the exact size, just relative comparison.
4. Once every item has an animal, translate animal → story point band for planning.

FACILITATION NOTE

Resist the urge to justify placements with numbers – the entire value of the exercise is training intuition, and numbers pull the room back into old habits.

WATCH OUT

Don't oversell this as "the Noah's Ark Game" with official rules – it's a themed skin on Animal Sizing, and being upfront about that keeps expectations honest.

The Paper Airplane Game

Teams fold as many quality paper airplanes as they can across timed rounds – but only one person may fold at a time, so hand-offs, bottlenecks, and work-in-progress piling up become impossible to ignore. Same iterative-improvement lesson as the Ball Point Game, with a flow/bottleneck angle the ball version doesn't teach as directly.

🕒 45–60 min 6–12 paper, a measuring tape

1. Explain the rule: one sheet of paper per plane, one fold per person, then pass it on – no one folds twice in a row on the same plane.
2. Team estimates how many finished planes they'll produce this round.
3. Run a 3-minute timed round.
4. A plane only counts as "done" if it's fully folded AND flies at least 3 metres when thrown – work-in-progress doesn't count, however close to finished.
5. Short retro, improve the process, run it again – five rounds total.

FACILITATION NOTE

The "must fly 3 metres to count" rule is what makes this teach Definition of Done, not just throughput – push teams to notice how many planes were "basically finished" but didn't actually qualify.

WATCH OUT

Needs a room with enough clear floor to actually throw the planes for the distance test – a small meeting room undersells the exercise.

The Group Counting Game

The whole team counts up from one, together, with no assigned order – anyone can say the next number, but if two people speak at once, the count restarts from one. Sounds trivial. Getting past ten as a group of eight is genuinely hard, and it's a sharp, fast lesson in reading the room instead of just waiting for your turn to talk.

🕒 10–15 min 5+ none

1. Explain the one rule: count up as a group, one number per person, no pre-set order – if two people say a number at the same time, go back to one.
2. Run it. Let the group fail and restart as many times as it takes.
3. Set a target (e.g. reach 20) rather than a time limit – this one ends on success, not the clock.

FACILITATION NOTE

For remote teams specifically, run this with cameras on and debrief the difference cameras make – reading someone's breath or posture before they speak is exactly the non-verbal signal remote work usually strips out, and this game makes losing it (or having it) obvious fast.

WATCH OUT

Don't rescue a stuck group by assigning an order – the discomfort of genuinely not knowing who goes next is the entire lesson.



TRY THIS AT HOME

Same technique, your actual life

Half of these work at the dinner table, verbatim. "One word for your day." "What's your battery at?" "Your week as a movie title." Kids especially take to Weather Report – "I'm foggy today" is emotional vocabulary disguised as a game. The Time Capsule is quietly one of the best New Year's Eve traditions a family or friend group can adopt.

AI SHORTCUT

The custom icebreaker generator

When the catalog feels stale, generate one matched to your actual week:

You're helping me open a team retrospective.
 Context: [team size]
 people, [remote/in-person/hybrid]. This sprint:
 [1–2 sentences –
 e.g. "shipped the big release but had two
 production incidents;
 team is tired but proud"]. Energy wanted: [calm/
 lively/silly].
 Time budget: [X] minutes.

Suggest 3 icebreaker options. For each: name,
 exact facilitator
 script (the words I say), time needed, and one
 way it could go
 wrong. Nothing that forces personal disclosure.
 Nothing that
 needs prep.

Also great for writing the Sprint Horoscope (card 25) – paste in three
 real sprint events and ask for mystic style. Thirty seconds of prep
 instead of five minutes.

Prefer a button to a prompt? Skip the prompt entirely – there's a whole
 tool built to do exactly this. See below.

Skip the copy-paste

Or let retro-generator.com run the whole retro for you

Every icebreaker, format and prompt in this book – generated, laid out and ready to run in seconds. No blank page, no wrestling an AI prompt into shape, no cobbling a board together five minutes before the call.

- **Tuned to your team.** Tell it your size, your mood and the sprint you just had – it builds the session around *that*, not a generic template.
- **The whole session, sorted.** Icebreaker, main activity and check-out picked and laid out for you, using the same mood-first logic as this chapter.
- **Design done for you.** A clean, ready-to-run board with zero formatting and zero blank-page dread.
- **Free to try.** You could spin up your next retro before you finish reading this chapter.

Try retro-generator.com →

Prefer a button to a prompt? This is the button.

FACILITATOR'S CHEAT SHEET

The whole chapter, in one box

- **The job:** every voice makes a sound in the first 5 minutes. Nothing more.
- **Match energy to need:** flat room → energy lifter; tight schedule → temperature check; polite-but-guarded room → trust builder.
- **Simultaneous beats sequential** – chat-storm, fist-to-five, emoji drop. No anchoring on the first answer.

- **Silence first:** a short thinking pause before anyone speaks noticeably improves answer quality – the classroom research on this is called "wait time" [Rowe, 1972], and the same mechanism holds in a retro.
- **Always allow a pass.** An icebreaker someone survived is worse than one they skipped.
- **You go first** when the icebreaker involves any vulnerability – and make it real.
- **Rotate.** Any icebreaker four retros in a row becomes furniture.
- **A stalled or heated room?** Icebreakers aren't just openers. A 90-second "would you rather" round (its own deck above) is a pattern-interrupt that resets energy *mid-retro*, not only at the start.
- **The answers are data.** Three disaster-movie titles and a 20% battery? Your retro agenda just wrote itself.
- **Don't just open well - close well too.** The Closing Check-Outs deck is the same trick pointed at the end of the meeting instead of the start.
- **Everything here fits a normal 45-min-1hr retro** except the Full Activities section – book real extra time for those.

Rapid-Fire Facilitation

Six things happen in the room. One move each, no overthinking - circle A or B, then check the key.

1

The timer hits zero mid-sentence, during the most honest moment of the retro.

☐ A

Cut it off - time's up and the next retro depends on you holding the clock.

☐ B

Extend consciously, say so out loud, and trim a later, less critical stage to make up the time.

2

The whole room is flat - one-word answers, phones face-up, going through the motions.

☐ A

Push harder into the agenda you planned.

☐ B

Name it and run a quick anonymous ESVP check to find out where the energy actually is.

3

You're asked to facilitate, but you're the Scrum Master and this retro is about the process you own.

- ☐ A Run it yourself - you own the calendar invite anyway.
- ☐ B Rotate the chair to someone with less riding on the answer.

4

You ask a question. Four seconds of silence. It's excruciating.

- ☐ A Jump in with your own answer to get things moving.
- ☐ B Say nothing, hold the room, and count to ten.

5

One person answers every question first, at length, and the rest have started deferring to them.

- ☐ A Tell them, in front of everyone, that they're dominating.
- ☐ B Switch to written, simultaneous input so volume can't own the format.

6

Someone raises a real point that's genuinely off today's topic, and the room starts following it.

A

Shut it down - it's not what we're here for.

B

Car Park it in front of everyone, so it's deferred and remembered, not dismissed.

Answer key

Best answer is **B** every time. Six for six and you're running the room, not the script.

1 · TIMEBOXING

B - extend consciously, trim later

Set visible timers so the room shares the constraint, and when something important is unfolding, extend out loud and cut a less critical stage instead. Cutting off the honest moment kills the value; ignoring the clock kills every retro after this one.

2 · WHOLE-ROOM DISENGAGEMENT

B - name it, run an ESVP check

A flat room is information, not a failure to push through. An anonymous Explorer/ Shopper/Vacationer/Prisoner read tells you where people actually are without spotlighting anyone - then shorten what's left or ask why.

3 · WHO FACILITATES

B - rotate the chair

When the retro judges the Scrum Master's own process, they shouldn't also run it - that's a conflict of position, not intent. Rotating also spreads a genuinely useful skill instead of hoarding it in one head.

4 · SITTING WITH SILENCE

B - wait, count to ten

The "wait time" finding [Rowe, 1972]: push the pause past a few seconds and answers get measurably longer and better. The silence is doing its job; don't rescue the room from it.

5 · THE OVER-TALKER

B - switch to written input

Redirect structurally rather than calling anyone out. Simultaneous written gathering means one voice can't dominate the format itself, and a public rebuke just costs that person their safety and derails the room.

6 · OFF-TRACK

B - Car Park it

A genuine but off-topic point goes into a visible Car Park so it's deferred, not dismissed - writing it down is what makes that credible. (For a conversation that's on-topic but circling, the tool is ELMO instead.)

PART III

The Retro Format Cookbook

30+ formats, recipe-card style

Pick Your Format: A Decision Guide

One page to kill the "we always do Mad/Sad/Glad" rut. Find your team's situation, follow the arrow.

The mistake most teams make: running the same format every single sprint because it's the one everyone remembers, regardless of whether it fits what actually happened. A format is a tool matched to a situation – a sailboat retro on the sprint after a production incident is the wrong tool, no matter how well you facilitate it.

Read down the list, find the row that matches your team this week, and go straight to that format. Chapters 6 through 9 have the full recipe cards.

Team is flat, low energy	Reach for a creative, high-energy format – Movie Poster Retro or Retro Bingo. Fun resets the room faster than more discussion does.
Same issues keep recurring	Don't reach for a new format – read Chapter 10 first. This is a follow-through problem, not a format problem.
Just shipped something big	Run the Post-Release format – celebrate before you dissect. Order matters.
Just had a production incident	Use the blameless Timeline format – reconstruct what happened before anyone assigns cause.
New team, or teams just merged	Start with a trust-builder icebreaker, then the gentle, structured Starfish format.
Team is remote or distributed	Go to Chapter 9 – formats built for remote, not classics awkwardly ported to a video call.
Team seems TOO comfortable	Run the Pre-Mortem – deliberately hunt for the risk nobody's naming.
Quarter or milestone check-in	Zoom out with the Quarterly Milestones Map instead of a normal sprint retro.

Heading into a break	
	Use the Pre-Holiday Wind-Down – close the chapter before everyone logs off.
Nothing special - a normal sprint	A classic done well is genuinely fine. Start/Stop/Continue or 4Ls, with the twist that keeps it fresh.

If two rows apply at once - say, remote AND just had an incident - solve for the sharper edge first. An incident retro run badly does more damage than a slightly generic remote adaptation.

FACILITATOR'S CHEAT SHEET

The one rule behind this whole chapter

— **Match the format to the sprint, not to habit.**

"We always do X" is a reason to change it, not a reason to keep it.

- **Recurring issues aren't a format problem.** No activity fixes a broken follow-through habit – go to Chapter 10 instead.
- **Celebration before surgery.** Post-release retros that skip straight to "what went wrong" waste the team's one chance to feel good about shipping.
- **When two situations collide, solve the sharper one first.**

The Right Retro for the Job

You're the retro coach on call this fortnight, and eight squads each ping you with their situation. One question for each: which format do you send them in with? Decide, then check the key.

How this works. Chapter 5 is the decision guide - find your team's situation, follow the arrow. This is the drill: eight different teams, eight situations, and the same job each time - match the format to the sprint, not to habit. Jot A, B, C or D for each, then turn to the answer key. The one rule underneath all of it: "we always do Mad/Sad/Glad" is a reason to change the format, not to keep it.

Eight squads, eight calls

Each team below tells you what their fortnight looked like. Pick the format that fits the situation - the right tool, not the most familiar or the most elaborate one.

QUESTION 1 · THE KESTREL SQUAD

"We had a production incident on Thursday - a bad deploy took checkout down for an hour. Feelings are still raw." What format do you send them in with?

- ☐ A Mad / Sad / Glad, to let people vent the emotion of it.
- ☐ B A blameless Timeline - reconstruct exactly what happened, step by step, before anyone assigns a cause.
- ☐ C A Sailboat, to look at what's pushing the team forward and holding it back.
- ☐ D Start / Stop / Continue - quick and familiar, just get it done.

QUESTION 2 · THE VELA TEAM

"We just shipped the launch we've been building for nine months. The whole channel is euphoric." What retro do you run?

- ☐ A Straight into "what went wrong" while the details are fresh - celebrate some other time.
- ☐ B A Post-Release retro that celebrates the win first, then dissects - order matters.
- ☐ C A Pre-Mortem on the next launch, to stay ahead of it.
- ☐ D A blameless Timeline of the whole launch week.

QUESTION 3 · THE ORCA SQUAD

"Our retro board has had the same three complaints on it for five sprints running." What do you reach for?

- ☐ **A** A brand-new, exciting format to shake things up.
- ☐ **B** Retro Bingo, to at least make it fun again.
- ☐ **C** Not a new format at all - this is a follow-through problem; fix how actions get owned and tracked (Chapter 10).
- ☐ **D** A Pre-Mortem, to dig into what's really going wrong.

QUESTION 4 · THE LUMEN TEAM

"Honestly the team's flat - one-word answers, phones on the table, going through the motions." Which format?

- ☐ **A** A serious, deep-dive root-cause session to force some engagement.
- ☐ **B** A creative, high-energy format - a Movie Poster Retro or Retro Bingo. Fun resets a flat room faster than more discussion.
- ☐ **C** Cancel it - there's no point running a retro when energy is this low.
- ☐ **D** The same classic as always, but push them harder to join in.

QUESTION 5 · THE ATLAS SQUAD

"Two teams just merged into one. Half of us have never worked together." How do you open their first retro?

- ☐ **A** Dive straight into Mad / Sad / Glad to get honesty fast.
- ☐ **B** Skip the icebreaker to save time - they'll warm up on their own.
- ☐ **C** A trust-builder icebreaker first, then the gentle, structured Starfish format.
- ☐ **D** A Pre-Mortem, so they bond over shared risks.

QUESTION 6 · THE CIRRUS TEAM

"Everything's fine, honestly. Velocity's steady, nobody's flagging any risks, it's all pretty smooth." Which format?

- ☐ **A** A celebratory Post-Release retro - they've earned it.
- ☐ **B** Start / Stop / Continue - keep the good times rolling.
- ☐ **C** A Pre-Mortem - deliberately hunt for the risk nobody's naming.
- ☐ **D** Mad / Sad / Glad, to double-check the mood.

QUESTION 7 · THE NIMBUS SQUAD

"We're fully remote across three timezones - and we had a production incident yesterday." Two situations collide and you can only optimise for one. Which edge do you solve first?

- ☐ **A** Pick a generic remote-friendly format and treat it as a normal sprint retro.
- ☐ **B** Run the blameless incident Timeline, adapted for remote - solve the sharper edge first, because a botched incident retro does more damage than a slightly generic remote one.
- ☐ **C** Postpone the retro until everyone can get in a room together.
- ☐ **D** Split it into two retros - one for remote issues, one for the incident.

QUESTION 8 · THE PIKE TEAM

"Completely normal sprint - nothing dramatic shipped, nothing broke, no drama at all." You feel like you ought to run something clever. What's the right call?

- ☐ **A** Invent an elaborate themed retro so it doesn't feel boring.
- ☐ **B** A classic done well - Start/Stop/Continue or 4Ls with a small twist to keep it fresh - is genuinely the right tool.
- ☐ **C** Skip the retro - nothing happened, so there's nothing to discuss.
- ☐ **D** A blameless Timeline, just to keep the format sharp.

Answer key

Give yourself a point for each best answer. **8** - format whisperer; you matched the tool to every sprint. **6-7** - good instincts; one or two reached for habit over fit. **4-5** - still defaulting; re-read the decision guide before next sprint. **0-3** - the Mad/Sad/Glad rut; Chapter 5 is one page, keep it open. Every answer links to the format's full recipe card.

QUESTION 1 · THE KESTREL SQUAD

Best answer: B

A blameless Timeline.

After a real incident you **reconstruct what happened before anyone assigns cause** - a shared, blameless sequence keeps the retro from turning into a hunt for who to blame. A Sailboat here is the wrong tool no matter how well you run it (C); Mad/Sad/Glad (A) airs the feelings but never builds the timeline that finds the cause; and Start/Stop/Continue (D) is far too shallow for something this consequential.

QUESTION 2 · THE VELA TEAM

Best answer: B

A Post-Release retro - celebrate, then dissect.

Celebration before surgery, and order matters. A euphoric team gets one chance to feel good about shipping; skipping straight to "what went wrong" (A) wastes it and teaches people that finishing is just the start of a blame list. Pre-Mortem (C) and a launch-week Timeline (D) both jump to problem-hunting before the win has been marked.

QUESTION 3 · THE ORCA SQUAD

Best answer: C

Not a format problem - fix follow-through (Chapter 10).

The same complaints for five sprints is a **broken follow-through habit, and no activity fixes that**. A new format (A), a fun one (B), or a Pre-Mortem (D) all just re-surface the same issues in nicer packaging - the fix is how actions get owned, tracked and revisited, not which game you play on the day.

QUESTION 4 · THE LUMEN TEAM

Best answer: B

A creative, high-energy format like Retro Bingo or a Movie Poster Retro.

For a flat room, **fun resets the energy faster than more discussion does**. A heavier deep-dive (A) or pushing harder on the usual classic (D) drives the flatness deeper; cancelling (C) throws away the signal that something's up. Lift the room first, then you can actually talk.

QUESTION 5 · THE ATLAS SQUAD

Best answer: C

A trust-builder icebreaker, then the gentle Starfish.

A brand-new or just-merged team has **no trust banked yet**, so you build a little first and then use a gentle, structured format. Diving into raw candor (A) before trust exists gets you silence or performance; skipping the icebreaker to save time (B) skips the very thing that makes the rest work; a Pre-Mortem (D) is too pointed for a room still learning to talk to each other.

QUESTION 6 · THE CIRRUS TEAM

Best answer: C

A Pre-Mortem.

A team that's *too* comfortable has one real danger - **the risk nobody's saying out loud** - and a Pre-Mortem goes hunting for it on purpose. A celebratory (A) or keep-it-rolling (B) format just reinforces the blind spot, and Mad/Sad/Glad (D) checks the mood you already know is fine while missing the risk you don't.

QUESTION 7 · THE NIMBUS SQUAD

Best answer: B

The blameless incident Timeline, adapted for remote.

When two situations collide, **solve the sharper edge first**. An incident retro run badly does far more damage than a slightly generic remote adaptation, so run the Timeline and just borrow the remote techniques from Chapter 9 to deliver it. Treating a fresh incident as a normal sprint (A) buries it; postponing (C) lets memories fade and blame harden; splitting it in two (D) fractures the one conversation the team most needs to have together.

QUESTION 8 · THE PIKE TEAM

Best answer: B

A classic done well, with a small twist.

A normal sprint doesn't need a gimmick - **a classic done well is genuinely the right tool**. Reaching for something elaborate every single time (A) is its own kind of rut; skipping the retro (C) forfeits the steady habit of reflection that keeps a good team good; and running an incident format with no incident (D) is a tool mismatched to the moment, the exact mistake this whole chapter is about.

The rule behind every answer

- **Match the format to the sprint, not to habit.**
"We always do X" is a reason to change it.
- **Recurring issues aren't a format problem** - that's follow-through (Chapter 10), and no activity fixes it.
- **Celebration before surgery.** A post-release retro that skips the win wastes the team's one chance to enjoy shipping.
- **Incident? Reconstruct before you assign cause** - the blameless Timeline, every time.
- **Too comfortable is a risk, not a rest** - a Pre-Mortem hunts the thing nobody's naming.
- **When two situations collide, solve the sharper edge first.**
- **Also in the chapter:** Quarterly Milestones Map for a milestone check-in, and the Pre-Holiday Wind-Down heading into a break.

The Classics, Done Well

Seven formats you've probably run before – the twist that keeps each one fresh, and the way it usually goes flat.

Classics became classics for a reason – they work.

They also go stale fastest, because "we always do this one" is how teams end up on autopilot. Every card below has a twist: a small change that makes a familiar format feel new without changing what it's actually good at.

The Seven Classics

recipe cards – steal the twist even if you keep the format

01

Mad / Sad / Glad

Sort the sprint into three emotional buckets. The most accessible format there is – everyone understands the categories instantly.

🕒 30–40 min any stickies, 3 columns

1. Silent write: one item per sticky, sorted Mad / Sad / Glad.
2. Post and cluster similar stickies together.
3. Dot-vote the top 2–3 clusters to discuss.
4. Discuss, then commit to one action.

THE TWIST

Add a fourth column: "Surprised." It catches the stuff that's neither good nor bad but genuinely unexpected – often the most useful signal in the whole board.

HOW IT GOES FLAT

Becomes a Complaint Circle (Chapter 1) when "Mad" fills up and "Glad" stays empty. If that happens two sprints running, that imbalance IS the retro topic.

Start / Stop / Continue

The most action-oriented classic – every input is already halfway to a decision, because it's framed as a behavior, not a feeling.

🕒 **30–40 min** **any** **stickies, 3 columns**

1. Silent write: things to Start doing, Stop doing, Continue doing.
2. Cluster and dot-vote.
3. For the top item in each column, agree the specific next step.

THE TWIST

Require every "Stop" item to come with a proposed replacement. "Stop doing X" without "start doing Y instead" often just creates a gap nobody planned for.

HOW IT GOES FLAT

"Continue" gets skipped entirely because it feels like the boring column. Don't let it – Continue items are your appreciation moment in disguise.

4Ls - Liked, Learned, Lacked, Longed For

Devised by Mary Gorman and Ellen Gottesdiener

[Gorman & Gottesdiener]. Digs deeper than Mad/Sad/Glad because "Learned" and "Longed For" pull out growth and aspiration, not just reaction.

🕒 35–45 min any stickies, 4 columns

1. Silent write across all four columns.
2. Read "Learned" items aloud first – they double as low-key knowledge sharing.
3. Cluster the rest, discuss top items.

THE TWIST


Read "Learned" out loud as a group before discussing anything else. It's the column most likely to contain something genuinely useful to the whole team, and it sets a constructive tone before the harder columns.

HOW IT GOES FLAT

"Longed For" turns into a wish list nobody owns. Cap it at one item carried into the Decide stage.

Sailboat

A metaphor retro: wind (what's pushing you forward), anchors (what's holding you back), rocks (risks ahead), island (the goal). Great for teams who respond better to a visual than a list.



🕒 40–50 min any a drawn sailboat scene

1. Draw (or template) a boat, anchors, rocks, sun, and an island on the horizon.
2. Silent write stickies onto each part of the scene.
3. Walk the board together, anchors first – they're usually the most actionable.
4. Pick one anchor to cut loose this sprint.

THE TWIST

Ask "what's ONE anchor we can cut loose THIS sprint?" – the metaphor is vivid enough that "cutting an anchor" becomes a real commitment people remember, not just a line item.

HOW IT GOES FLAT

People spend all their energy being clever about the metaphor instead of being honest about the anchors. Keep redirecting: "so concretely, what is this anchor?"

Starfish

Devised by Patrick Kua [Kua, 2006]. Five-point spread: Keep Doing, Less Of, More Of, Stop Doing, Start Doing. More detailed than Start/Stop/Continue, which makes it gentler – nothing has to be framed as binary.

🕒 35–45 min any a 5-point star template

1. Silent write across all five points.
2. Cluster similar items within each point.
3. Discuss "More Of" and "Less Of" first – they're the least confrontational entry point.

THE TWIST

This is the best classic for a new or low-trust team (see Chapter 5) – "less of" feels much safer to write than "stop," because it doesn't demand a team admit something is entirely wrong.

HOW IT GOES FLAT

Five columns is a lot to fill honestly in one sitting. If energy is low, cut to three points rather than rushing all five.

Timeline

Reconstruct the sprint (or incident) chronologically before discussing anything. Removes the recency bias where only the last three days get remembered.

🕒 **40–60 min** **any** **a horizontal timeline**

1. Draw a horizontal line spanning the sprint (or incident window).
2. Everyone adds events, decisions, and mood markers where they happened in time.
3. Walk the timeline together, left to right, before any "why" discussion.
4. THEN generate insights and decide.

THE TWIST

Add a second row underneath purely for mood – a line graph of how the team felt at each point. Seeing the emotional dip line up exactly with a specific Tuesday deploy is more persuasive than anyone describing it in words.

HOW IT GOES FLAT

Skipping straight to discussion before the timeline is fully built.

Reconstruction and analysis are separate stages – don't let them blur.

Lean Coffee

Created by Jim Benson and Jeremy Lightsmith [Benson & Lightsmith] as a meetup format, since adopted widely for retros. No pre-set topics at all – the team builds the agenda live, in the room, and votes on what actually matters to discuss right now.

🕒 30–45 min 4–10 stickies + a timer

1. Everyone silently writes discussion topics, one per sticky.
2. Dot-vote to rank them.
3. Discuss the top topic for a timeboxed 5 minutes.
4. At the buzzer: thumbs up (keep going), down (move on), sideways (one more minute). Repeat down the ranked list.

THE TWIST

Use the thumbs vote strictly – it's what stops one topic eating the whole retro, and it hands the room real-time control over pacing instead of the facilitator guessing.

HOW IT GOES FLAT

Works best for teams who already trust each other enough to vote honestly. On a brand-new team, the ranked topics tend to be safe ones – pair with a trust-builder icebreaker first.

FACILITATOR'S CHEAT SHEET

Seven classics, seven twists

- **Mad/Sad/Glad:** add a "Surprised" column for the signal the other three miss.
- **Start/Stop/Continue:** every Stop needs a proposed Start in its place.
- **4Ls:** read "Learned" aloud first – free knowledge sharing.
- **Sailboat:** commit to cutting one anchor loose THIS sprint.
- **Starfish:** the gentlest classic – best for new or low-trust teams.
- **Timeline:** reconstruct chronology fully before any "why" discussion starts.
- **Lean Coffee:** the thumbs vote is what stops one topic eating the whole retro.

Formats for Specific Moments

Six situations that need something other than your default retro – and exactly what to run instead.

A generic format on a non-generic sprint wastes the moment. The sprint after an outage, the sprint after a launch, and the sprint before a holiday are not "just another sprint" – each deserves a format shaped for what actually happened.

Six Specific Moments

matched formats for when "normal" doesn't fit

01

Post-Incident: The Blameless Timeline

Something broke in production. Feelings are high, blame is one careless sentence away, and the facts of what happened are still scattered across five people's memories.

🕒 **60–90 min** **responders + team** **timeline template**

1. Read the Prime Directive – louder and more deliberately than usual.
2. Build the timeline together, facts only: what happened, when, who noticed, what was done. No "why" yet.
3. Once the timeline is agreed, ask "why" of the process, never of a person: "what about our systems let this happen?"
4. Identify contributing factors (plural – incidents are rarely one cause) and 1–2 concrete follow-ups.

FACILITATION NOTE

Ban the phrase "you should have" outright, out loud, at the start. If it slips out anyway, redirect immediately and calmly: "let's rephrase that as a system question."

WATCH OUT

Don't run this the same day as the incident if people are still in fight-or-flight mode – 24–48 hours later gets you facts without raw adrenaline.

Post-Release: Celebration + Hangover

You just shipped something big. The instinct is to jump straight into "what went wrong" – resist it. Order matters here more than in any other format.

🕒 45–60 min any none, maybe drinks

1. Open with the Launch Day Toast icebreaker (Chapter 4, card 23).
2. Celebration round: everyone names one proud moment from the release, no "buts" allowed.
3. THEN the hangover round: what was harder than it should've been, gather as normal.
4. One action item for next time – same as any retro.

FACILITATION NOTE

Physically separate the two rounds – different boards, a short break between them. Negative information reliably carries more psychological weight than positive information of the same size [Baumeister et al., 2001] – on a shared board, a handful of criticisms will crowd out a room full of wins unless you protect the celebration round with real separation.

WATCH OUT

Someone always tries to sneak a "but" into the celebration round. Gently stop them: "hold that thought for the next section."

Quarterly Zoom-Out: The Milestones Map

Sprint retros are myopic by design – that's their strength. Once a quarter, zoom all the way out and look at the shape of the whole period instead.

🕒 90 min whole team a long horizontal timeline

1. Build a timeline spanning the whole quarter – releases, incidents, team changes, big decisions.
2. Mark energy/mood across the same span (borrow the Timeline classic's mood row, Chapter 6).
3. Look for the shape: was this a quarter of steady progress, a slow grind, a sprint-then-crash pattern?
4. Set ONE theme for next quarter, not a list of ten fixes.

FACILITATION NOTE

Pull real data in beforehand if you can – deploy counts, incident counts, velocity trend – so the room's memory isn't the only input (this is exactly the kind of prep Chapter 12's AI shortcuts are built for).

WATCH OUT

Ninety minutes is a lot to ask for. Protect it on the calendar weeks ahead, and treat a cancellation as a real loss, not a free pass.

Pre-Holiday Wind-Down

The last working day before a break. Nobody's focused on next sprint's blockers – and forcing that focus wastes everyone's goodwill on the way out the door.

🕒 30 min any none

1. Run the Holiday Special icebreaker (Chapter 4, card 22).
2. Quick round: one thing from this period you're proud of.
3. Skip the usual "gather data / decide" stages entirely – this isn't a normal retro.
4. Close early. Let people leave with time back, not less.

FACILITATION NOTE

This is the one format in the book where ending early is a feature, not a failure. Thirty minutes, tops.

WATCH OUT

Don't sneak a real retrospective in under the wind-down label – the team will notice, and it undermines the trust you're building elsewhere.

The Team Is Grumpy: Energy Audit

Morale has been low for a while and nobody's named why. A normal retro format risks becoming another Complaint Circle (Chapter 1) – this one is built to find the actual source.

🕒 45 min any a simple energy-over-time chart

1. Everyone draws their own personal energy line over the last month – a rough sketch, not data.
2. Share the shapes: who's declining, who's flat, who's actually fine?
3. For each dip, ask what caused it – work, process, something outside work entirely (and it's fine if people don't share that part).
4. Identify ONE systemic factor the team can actually change.

FACILITATION NOTE

Some dips will be personal, not team-related – don't force everyone's line to have a team-fixable cause. The goal is finding the SHARED dip, if one exists.

WATCH OUT

If everyone's line has been low for months, this single retro isn't enough – escalate to leadership rather than expecting the format to fix a sustained burnout pattern.

The Team Is TOO Comfortable: Pre-Mortem

Everything feels fine. Suspiciously fine. Sometimes a smooth run means nobody's surfacing the risk quietly building underneath. The technique itself comes from decision researcher Gary Klein [Klein, 2007] – imagining a failure has already happened turns out to surface real risks that a straightforward "any concerns?" rarely does.

🕒 45 min any none

1. Frame it: "It's six months from now and this project has failed badly. What happened?"
2. Silent write: everyone's version of the failure story.
3. Cluster the stories – recurring failure themes are your real risks.
4. Pick the most plausible one and ask: what would we need to see NOW to catch this early?

FACILITATION NOTE

Assign a rotating "devil's advocate" whose explicit job is to disagree with the room's comfort – it depersonalizes dissent, since it's a role, not a personality trait.

WATCH OUT

Don't run this so often it manufactures anxiety in a genuinely healthy team. Once a quarter is plenty unless something specific is nagging at you.

Six moments, at a glance

- **Post-incident:** facts before "why," and never a "why" aimed at a person.
- **Post-release:** celebrate fully before you critique – physically separate the two.
- **Quarterly zoom-out:** one theme for next quarter beats ten fixes for this one.
- **Pre-holiday:** thirty minutes, ends early on purpose.
- **Grumpy team:** hunt for the shared dip, not every individual's personal one.
- **Too comfortable:** a rotating devil's advocate role depersonalizes dissent.

Creative & High-Energy Formats

Permission to be silly – and why silly gets more honesty out of a tired room than serious does.

Counterintuitive but true: a room laughing about a ridiculous metaphor will tell you more honest things than a room being asked earnestly to reflect. Play lowers the stakes just enough that people say the true thing instead of the careful thing. Use these when energy is flat, the last three retros felt the same, or the team just needs to remember work can be fun.

Don't underestimate a good meme or a well-timed joke here, either – teams love laughing at their own expense, or at agile's expense (there is no shortage of material). If the perfect one doesn't already exist, that's not a reason to skip it: a quick meme generator takes less time than writing the caption in your head. A well-placed inside joke about last sprint's chaos does the same job as any icebreaker in this book – it just does it in one image instead of a script.

Seven High-Energy Formats

use sparingly – novelty is the whole point, so don't run the same one every month

01

Movie Poster Retro

Design a movie poster for the sprint – title, tagline, genre, star rating, one-line review. Turns a plain summary into something people actually enjoy making.

🕒 35–45 min any paper or a shared canvas

1. In small groups, sketch a poster: title, tagline, genre (comedy? horror? heist film?), lead "actors" (real teammates, in character).
2. Each group pitches their poster to the room in 60 seconds.
3. Debrief: what does the genre you picked say about how the sprint actually felt?

FACILITATION NOTE

The debrief question is where the real retro happens – the poster is just the vehicle to get people talking freely first.

WATCH OUT

Groups of 3–4 work better than solo – it becomes a collaborative bit rather than a performance a shy person has to carry.

Superhero Retro

Everyone picks (or is assigned) a superhero whose powers match something they brought this sprint – then names their kryptonite too.

🕒 30–40 min any none

1. Everyone picks a hero whose power matches a strength they used this sprint.
2. Share the pick and why, one at a time.
3. Then: "what's your kryptonite right now?" – the honest, harder half.
4. As a group, identify one kryptonite that's shared across multiple people – that's systemic, not personal.

FACILITATION NOTE

This is the extended-family version of the Superpower/Kryptonite icebreaker (Chapter 4, card 10) – run the icebreaker first if the team hasn't tried the concept before.

WATCH OUT

Skip this one with a team that finds comic-book framing cringey rather than fun – read the room before you pick a creative format at all.

Kitchen Nightmares

Facilitator plays a Gordon-Ramsay-style inspector reviewing the sprint like a failing restaurant kitchen. Exaggerated bluntness, played for laughs, that somehow makes real critique land softer.

🕒 35–45 min any none, some acting

1. Facilitator opens in character: "Right, what happened in THIS kitchen this sprint?"
2. Team pitches "dishes" (deliverables) that came out great and ones that came out raw.
3. Drop the character for the real discussion: what actually needs to change in the kitchen (process)?

FACILITATION NOTE


The bit works because it's clearly a performance, not a real accusation – that theatrical distance is exactly what makes blunt feedback feel safe to hear and give.

WATCH OUT

Needs a facilitator who's genuinely comfortable hamming it up. Forced, half-hearted theater is worse than skipping the bit entirely.

Weather Map

The team draws a literal weather map of the sprint – storms over the rough patches, sunshine over the wins, fog over the unclear bits. Visual thinkers love this one.



🕒 35–45 min any a blank map/board

1. Draw or template a blank "map" of the sprint (a timeline works as the base).
2. Everyone adds weather symbols where events happened – sun, storm, fog, rainbow.
3. Walk the map as a group, storms first.
4. For the biggest storm, ask: what would have kept us out of it?

FACILITATION NOTE

Fog symbols are the most interesting ones to dig into – they mark moments where nobody was sure what was happening, which is often a communication gap worth naming.

WATCH OUT

Very similar in mechanics to Timeline and Sailboat (Chapter 6) – rotate between them rather than running all three in a month.

Marketplace / Open Space

No single topic – instead, several "stalls" run in parallel, each hosted by whoever proposed the topic, and people move freely between them. Best for teams with genuinely varied concerns nobody wants forced into one conversation.

🕒 **45–60 min** **8+** **multiple boards/tables**

1. Open call: who wants to host a topic? Set up one "stall" per host.
2. Everyone else circulates freely for a set time, joining whichever stalls interest them.
3. Each host summarizes their stall's key point to the full room at the end.
4. Group votes on which single item becomes THE action item.

FACILITATION NOTE

This only works with genuine topic diversity and a big enough group to split meaningfully – under 8 people, just run Lean Coffee instead (Chapter 6).

WATCH OUT

Popular hosts can end up with a crowd while others sit alone. If that happens, gently redirect a few people: "mind splitting over to the quieter stall for five minutes?"

Retro Bingo

A bingo card of sprint clichés – "someone says 'it worked on my machine,'" "a meeting could've been an email," "the deploy breaks on a Friday." Pure levity, built to reset a genuinely exhausted room.

🕒 15–20 min, then normal retro any a bingo card

1. Before the retro, build a 5×5 card of running team in-jokes and clichés.
2. Hand it out, let people mark off what actually happened this sprint.
3. Whoever gets a line shouts "retro bingo!" – small prize optional.
4. Roll straight into a normal retro afterward – this is a warm-up, not a replacement.

FACILITATION NOTE

Crowdsource the card from the team itself over the week beforehand – in-jokes the team wrote are funnier than ones the facilitator guessed.

WATCH OUT

This is purely Chapter 2's "Set the Stage," not a full retro on its own – always follow it with the real gather/insights/decide stages.

The Gratitude Retro

An entire retro built around appreciation instead of improvement – for the sprint after something genuinely hard, when the team needs to feel good before it can look critically at anything.

🕒 30 min any none

1. Run the Appreciation Shower icebreaker (Chapter 4, card 12) as the whole retro, not just the opener.
2. Extend it: everyone also names one thing they're grateful for about the WORK itself, not just teammates.
3. Close with a single light action item if one genuinely surfaces – don't force one.

FACILITATION NOTE

This is the one format in the chapter where it's fine to skip the Decide stage from Chapter 2 entirely. Some sprints, the retro's whole job is helping the team feel okay again.

WATCH OUT

Don't run this as an avoidance tactic when the team actually needs the Post-Incident or Grumpy Team formats (Chapter 7) instead. Gratitude is a break, not a substitute for a hard conversation that's overdue.

Seven creative formats, one caution

- **Play lowers stakes** – a laughing room tells you the true thing more often than a serious one does.
- **Movie Poster / Superhero / Kitchen Nightmares:** the fun is the vehicle; the debrief question is the retro.
- **Weather Map / Marketplace:** best for visual thinkers and genuinely varied concerns, respectively.
- **Retro Bingo** is a warm-up, not a full retro – always follow with real gather/insights/decide.
- **Gratitude Retro** is a valid full retro after something hard – but never a dodge for a conversation that's overdue.
- **Rotate.** Novelty is the entire mechanism – repeat any of these monthly and it stops working.

Remote & Async Formats

Built for distributed teams from the ground up – not classics awkwardly ported onto a video call.

The usual failure: take an in-person format built around physical stickies and a shared wall, cram it into a video call, and wonder why nobody talks and the whiteboard tool crashes. These formats assume screens, time zones, and camera fatigue from the start.

Five Remote-Native Formats

designed for distributed teams, not adapted for them

01

Async Pre-Fill + Live Discussion

Splits the retro into two parts across time: written input happens async over 24 hours, live time is spent entirely on discussion. Solves the single biggest remote problem – people typing in real time is slower and less honest than people typing when they're actually ready.

🕒 **24h async + 30 min live** **any** **shared async board**

1. Open a shared board 24 hours before the live retro, one prompt per column.
2. Everyone contributes on their own time – no live pressure, no time-zone disadvantage.
3. Facilitator clusters overnight before the live call.
4. Live call is 100% discussion and deciding – reading and writing already happened.

FACILITATION NOTE

This format alone can cut live retro time by a third, because the slowest stage (Gather Data) already happened before anyone joined the call.

WATCH OUT

The board will be empty at hour 20 if you don't nudge it. Seed it yourself and post a reminder in the team channel with 4 hours left.

The 25-Minute Remote Retro

Remote meeting fatigue is real – a 60-minute video retro loses the room by minute 40. This is the full five-stage recipe compressed hard, for weeks when a full retro isn't realistic but skipping entirely isn't either.

🕒 25 min any a simple 3-column board

1. 0:00 – 1-minute check-in (Emoji Check-In or Fist to Five, Chapter 4).
2. 0:03 – 5-minute silent write, one board, no columns beyond Went Well / Went Badly.
3. 0:08 – 10-minute discussion of the top 2 clusters only.
4. 0:20 – 3-minute decide: one action, one owner.
5. 0:24 – close.

FACILITATION NOTE

Discipline is everything here – the timeboxes are tighter than feels comfortable at first. Visible timers are non-negotiable (see Chapter 3's timeboxing tip).

WATCH OUT

This shouldn't become the permanent default – it drops the depth a full retro gets. Use it for busy weeks, not every week.

Silent Writing Rounds

The entire "gather data" and part of "generate insights" happen with cameras on but nobody talking – just typing into a shared doc simultaneously, live. Removes the awkward unmuting dance remote meetings are full of, and sidesteps a documented problem with spoken group brainstorming called production blocking: only one person can talk at a time, so everyone else is either waiting or forgetting their own point [Diehl & Stroebe, 1987]. Writing simultaneously removes that bottleneck entirely.

🕒 10–15 min block any shared live doc

1. Open a shared doc, cameras on, mics off.
2. Everyone types simultaneously for a set window – no talking, no reacting yet.
3. Facilitator calls time, then opens discussion of what's on the page.

FACILITATION NOTE

The silence is the feature, not a glitch – resist any urge to narrate or fill the quiet while people type. Same principle as Chapter 3's "sitting with silence," just typed instead of spoken.

WATCH OUT

Doesn't work well over a shaky connection where the doc lags – test the tool beforehand if anyone's on unreliable wifi.

Tool-Agnostic Whiteboard Patterns

Not a single format – a set of layout patterns that transfer cleanly to any online whiteboard tool (Miro, FigJam, Mural, a plain shared doc) without needing the specific tool's fancy features.

🕒 **varies** **any** **any** **shared canvas**

1. Pre-build column/quadrant templates before the retro starts – never build structure live while people wait.
2. Color-code sticky types (issue = one color, idea = another) so patterns are visible before anyone reads a word.
3. Lock a "parking lot" area for off-topic points so the main board stays clean – see Chapter 3 for Car Park and its counterpart, ELMO, for circular conversations.
4. Export or screenshot the board at the end – remote retros lose their artifact if nobody saves it.

FACILITATION NOTE

Send a 60-second "how to use this board" video or GIF before the retro for anyone unfamiliar with the tool – troubleshooting software live burns the same goodwill an icebreaker just built. For specifics on Miro vs. Mural and building a reusable shape library, see Your Virtual Toolkit in the back matter.

WATCH OUT

Don't chase every fancy feature a whiteboard tool offers. The simplest board everyone can actually use beats the cleverest one two people can't navigate.

Breakout Rotation

For larger remote teams (10+) where one big call means several people never speak. Small breakout rooms discuss in parallel, then rotate or report back.

🕒 40–50 min 10+ video breakout rooms

1. Split into groups of 3–4 in breakout rooms, same prompt for everyone.
2. 10 minutes of discussion per group – small enough that everyone actually talks.
3. Back in the main room: each group reports their single most important point.
4. Full group votes on which reported point becomes the action item.

FACILITATION NOTE

Assign a scribe per breakout room in advance – without one, the report-back is always vaguer than the actual conversation was.

WATCH OUT

Breakout rooms can feel disconnected from "the real retro." Always close with everyone back together, never end while still split up.

FACILITATOR'S CHEAT SHEET

Five remote formats, one principle

- **Split time and talk.** Async pre-fill moves the slowest stage off the clock entirely.
- **Compress deliberately** – the 25-minute retro is for busy weeks, not every week.
- **Silent typing beats the unmuting dance** – let quiet be the format, not an accident.
- **Build structure before the call starts.** Never construct a template live while people wait.
- **Large team? Go small first.** Breakout rooms of 3–4, always reconvened before closing.
- **Save the artifact.** A remote retro with nothing exported afterward might as well not have happened.

From Talk to Change: Action Items That Actually Happen

The difference between a fun meeting and a team that actually gets better – and what to do when nothing ever changes.

Here's the uncomfortable truth: a retro can hit every stage from Chapter 2, use a brilliant format from this cookbook, get real laughs from a Chapter 8 icebreaker, and still fail completely – if nothing said in the room ever actually changes. Fun is not the goal. Better is the goal. Fun is just how you get there.

The graveyard, and how it forms

Go back to Chapter 1's Action-Item Graveyard: six good ideas, no owners, nothing done. It doesn't happen because teams are lazy – it happens because ambition without structure defaults to nothing. Six unowned commitments are, in practice, zero commitments. Everyone assumes someone else has it.

One action item, one owner, one check-in. That's the entire system. Everything else in this chapter is just making that system stick.

Rule one: pick ONE thing

Every retro will surface more good ideas than the team can act on. That's fine – write them all down somewhere, a backlog of improvements. But the retro itself commits to exactly one. Not the biggest one, not the most urgent one necessarily – the one most likely to actually get done before the next retro. Momentum beats ambition.

Rule two: name an owner, out loud, in the room

"The team will improve documentation" has no owner. "Priya will draft a documentation template by Friday and share it in standup" has one. The gap between those two sentences is the entire gap between change and the graveyard.

SAY THIS AT THE DECIDE STAGE

"Okay – [name], are you good to own this? What does 'done' look like, and by when?" Get a real answer to all three parts before moving on. A nod isn't an answer.

Rule three: make it visible between retros

An action item that lives only in the retro notes doc is an action item nobody sees again until the next retro reminds them it never happened. Put it somewhere the team already looks – pinned in the team channel, a card on

the sprint board, a line in standup notes. Visibility is what turns a good intention into an ambient reminder.

Rule four: open every retro by checking the last one

This is the single most useful habit in this entire book. Ninety seconds, every retro, before anything else: "last time we committed to X, owned by Y – did it happen?" Three outcomes, and all three are useful:

- **It happened.** Say so out loud. Momentum is contagious, and the room needs to see that this process actually works.
- **It's in progress.** Fine – check the timeline is still realistic, move on.
- **It didn't happen.** Don't skip past this one. Ask why, honestly, without blame (Prime Directive again). Sometimes the answer is "it wasn't actually the right priority" – that's useful information too.

When nothing ever changes

Sometimes a team follows every rule above and action items still quietly die, sprint after sprint. When that happens, stop running normal retros for a cycle and make follow-through itself the retro topic. Ask directly: what's actually stopping us from doing the thing we agreed to do? Common answers, and what they usually mean:

"We didn't have time"

Usually means the action item competed with sprint work and lost every time, because nobody protected space for it.

WHAT IT MEANS

The commitment needs to be small enough to fit in the cracks, or it needs explicit calendar time – not just good intentions.

"It wasn't really up to me"

The named owner didn't actually have the authority or access to make the change happen.

WHAT IT MEANS

Check ownership against actual authority before the retro ends, not after it's already failed.

"Honestly, I forgot"

The visibility rule (three) got skipped – the commitment lived only in a doc nobody reopened.

WHAT IT MEANS

Fix the reminder system, not the person. This is a process gap, not a character flaw.

"It didn't seem that important anymore"

Priorities shifted after the retro and the action item quietly lost its urgency.

WHAT IT MEANS

That's a legitimate outcome – but say so explicitly in the room instead of letting it silently vanish. A named "we're dropping this" beats an unspoken one.

TRY THIS AT HOME

The one-thing rule, at home

Same trap happens with New Year's resolutions and family "we should really – " conversations: five good ideas, zero owners, zero follow-up. Pick one change, name who's doing it (even if it's just you), and set a specific date to check whether it actually happened. It's the exact same fix, just without the sprint board.

FACILITATOR'S CHEAT SHEET

Four rules that close the graveyard

- **One action item, not six.** Momentum beats ambition – write the rest in a backlog, commit to just one.

- **Named owner, said out loud, with a "done" definition and a date.** A nod isn't an answer.
- **Make it visible** somewhere the team already looks – not buried in a notes doc.
- **Open every retro by checking the last commitment.** Ninety seconds, most useful habit in the book.
- **If it keeps failing, make follow-through itself the retro topic** – the four common excuses above each point to a specific, fixable gap.

PART IV

Your AI Co-Pilot

*AI handles the paperwork, humans handle
the talking*

The Retro That Almost Wasn't

A short story.

The rest of this book is a field guide. This chapter isn't – it's a short story, and it sits here because there's a difference between being *told* what AI can do for a retro and watching a team discover it. The practical chapters start again at Chapter 11.

Monday

Mark Sullivan had stopped believing in retrospectives six sprints ago.

Not officially. Officially, he still stood at the whiteboard every second Friday with his marker capped and ready, still drew the three columns – *Went Well, Didn't Go Well, Actions* – still asked the team to "be honest, this is a safe space." But he knew, and he suspected everyone in the room knew, that the same five sticky notes would go up in the same three columns they always did. *Not enough time for testing. Too many meetings. Communication with the platform team.* Someone would half-heartedly own an action. Nobody would look at it again until the next retro, when it would reappear, a little more sun-bleached, a little more resigned.

Meridian Insurance's claims-processing platform was eleven days from a regulatory deadline, and the Atlas squad – Mark's squad – was the one holding the rope. Three sprints of slipped commitments. A director, two levels above him, had started attending stand-ups. Not to help. To watch.

"We need to talk about the actions from last retro," Mark said, sixty seconds into the ceremony, already knowing how this would go.

Dev lead Priya didn't look up from her laptop. "Which ones? We had nine."

"That's kind of the problem," said Tom, the most senior engineer on the team, arms crossed. "Every retro we generate a graveyard of actions nobody does. Can we just skip this one? We're eleven days out. I've got failing integration tests calling my name."

Nobody objected. That was the part that stung. Not the pushback – the *silence* afterward. Six people who, a year ago, would have argued the point either way, now just waiting for permission to leave.

Mark looked at the whiteboard, at the ghost of a hundred retros layered underneath the current one, and uncapped his marker anyway.

"Fifteen minutes," he said. "Give me fifteen minutes."

He got twelve, and got nothing from them worth the twelve.

The Mentor

He found Desmond Achterberg in the building's least-used stairwell, which was where he always was around 4 p.m., supposedly getting steps in, actually avoiding his own calendar. Desmond had run agile transformations at two other companies before Meridian hired him as a principal coach, and he had a habit of answering questions with questions that made you feel briefly stupid and then, a day later, considerably less so.

"Your retros are theater," he said, before Mark had finished explaining. "You know that."

"I know the *actions* aren't landing. That's different from theater."

"Is it? Theater is a group of people performing a ritual whose outcome is known in advance." Desmond stopped climbing. "When's the last time a retro at Meridian surfaced something nobody in the room already knew?"

Mark opened his mouth, then closed it.

"That's not a knock on you," Desmond said. "It's a knock on the format. Sticky notes and a fifteen-minute time-box were designed for a team that meets in one room, ships every two weeks, and has nothing at stake but its own velocity. Your team is distributed across three time zones, shipping into a regulated platform, carrying eleven months of unresolved friction with a platform team that doesn't report to you. You're trying to excavate a year of scar tissue with a Sharpie."

"So what – I need a bigger Sharpie?"

"You need memory," Desmond said. "And you need someone in that room willing to say the thing that gets six people uncomfortably quiet, without it costing them anything to say it."

He didn't finish the thought. Desmond rarely did. He just handed Mark a name – Rhea Vance, an engineering manager two floors down whose team had quietly stopped struggling about four months back – and went back to his stairwell.

What Rhea's Team Had Done

Rhea's team's retro board didn't look like sticky notes. It looked like a living document that remembered.

"We stopped running retros from scratch every time," Rhea explained, pulling it up on the shared screen. "Every retro references the last six. Every action either got done, got explicitly killed, or is still open – and it *says so*,

out loud, before anyone writes a new sticky note. You cannot add a tenth open action to this board without addressing the nine that are already sitting there."

"That sounds like it would make people defensive."

"It did, for one retro. Then it made them relieved, because for the first time the pattern was visible to *them*, not just felt by them. We started feeding the raw retro notes – anonymized – into a tool that clusters them across sprints. It doesn't tell us what to think. It tells us: 'this exact complaint, worded four different ways, has come up in five of your last seven retros, and every time the action was assigned to someone who doesn't control the root cause.' That's not a discussion topic. That's a diagnosis."

Mark frowned. "You're using AI to run your retros?"

"I'm using it to hold up a mirror we're too tired to hold up ourselves. The facilitation is still human. The courage to say the hard thing out loud is still human. What the tool does is refuse to let the room forget. It doesn't care about politics. It doesn't care that Tom is exhausted and just wants to get back to his failing tests. It just says: here is the pattern, here is how long it's been true, here is who it actually depends on to fix."

She showed Mark a report from three retros back. Buried in eleven paraphrased comments about "communication with the platform team" was a line the clustering tool had pulled to the top: four separate people, across three sprints, describing the same handoff failure in different words – a schema change that platform made without notice, twice, each time costing Rhea's team two days of rework.

"Nobody ever put that in one sentence before," Rhea said. "We were all describing the elephant from a different side of the room."

What Mark Overheard

He didn't go looking for it. That was the thing he'd remember later – he wasn't hunting for a theme, he was just in the room, the way he was in the room every day.

Monday stand-up, Priya update: "I feel like I got voted off this sprint and nobody warned me first." A throwaway line about the platform team's schema change. Mark almost didn't register it – until Tom, without looking up from his screen, muttered, "classic blindside," and two people laughed the specific kind of laugh that meant *everyone gets that reference except me if I don't watch the show*.

He let it go. Coincidence.

Tuesday, at the coffee machine, he caught the tail end of a conversation between Femi and one of the QA engineers. " – no because he was playing both alliances the whole time, that's why nobody trusted him at the vote – " Femi stopped when he saw Mark, half-embarrassed, like he'd been caught slacking. He hadn't. It was 8:52 in the morning. He'd just been talking about the show.

Wednesday, a Slack thread he was only half paying attention to – the one where the team dumped memes and complaints in equal measure – had six messages in a row about who was going to get "voted off the island" this week, in a tone that was clearly not about the show anymore. It was about the sprint. About who felt exposed. About who'd been left holding a ticket nobody backed them up on.

That was the third time. Mark sat back from his laptop.

He'd been so focused on the *events* of the sprint – the hotfixes, the provisioning delay, Femi's first release – that he'd almost missed this entirely, because it wasn't happening in a stand-up update or a Jira ticket. It was happening in the gaps: the two minutes before a call started, a Slack thread that wasn't about work at all. Nobody puts "obsessed with a survival show" on a sprint report. But it was arguably telling him more about how to reach this team than the report was.

Blindsided. Voted off. Alliances. Nobody warned me. They weren't just watching a survival show together. They were narrating the sprint through it, without noticing they were doing it – because it was easier to say "I got blindsided" out loud than "I feel like this team doesn't have my back," even though everyone in the room would have heard exactly that.

He thought about the deadline, the midnight hotfixes, Femi's first prod release, the demo that had gone better than anyone expected, the two days lost to provisioning nobody had flagged until it was too late. That was the sprint's story. Then, separately, he thought about how flat the room had felt lately – not the *what*, just the *how*. Not broken, but wary, like everyone was quietly keeping score of who'd let them down and who hadn't, and nobody was saying it out loud. That wasn't something anyone had told him. It was just how the last two stand-ups had felt to sit in.

That was the click. Not a checklist. Just three things he'd been half-hearing all week suddenly sitting next to each other and making one sentence: *this team already has the language for what they need to say, they just don't know they're using it.*

He didn't have the energy, at 4 p.m. on a Wednesday with a deadline breathing down his neck, to invent a themed retro from scratch. So he didn't. He opened the same assistant he'd used for the pattern summary and typed out, roughly: *my team keeps making survival-show references about feeling blindsided and not trusting who's got their back – we just shipped under pressure, had a rough provisioning delay, and the mood is wary but not broken. What retro format would fit that?*

It came back with three options in under a minute. The one that stopped his scrolling was built entirely out of the show's own vocabulary – *Safe This Round* instead of *Went Well*, *Blindsided By* instead of *Didn't Go Well*, and a closing round called *Who'd You Vote To Keep*, for whatever process was actually worth protecting going into the next sprint.

He hadn't invented anything. He'd just finally listened to what the team had been telling him all week, and let something else do the five minutes of work turning it into a session he could actually run.

Wednesday, Nine Days Out

Mark walked into the next scheduled retro and framed it differently before anyone had even sat down: "Thirty minutes. I want to try something. If it's useless, we go back to the old way and I'll stop pushing."

Under *Blindsided By*, he fed the last five retros' raw notes – with names stripped – into the same clustering approach Rhea used. He didn't tell the team what it found. He let the tool's summary go up on the screen cold.

Pattern detected across 5 retrospectives: "insufficient time for testing" appears in 7 forms. In every instance, the described root cause traces to environment provisioning delays from the platform team, not testing capacity. Related action items (4) were assigned to Atlas squad members with no authority over provisioning.

The room went quiet in a different way than usual.

"That's – " Priya started, then stopped, rereading it. "We've been fixing the wrong thing for four sprints."

"We've been fixing the *only* thing we had authority to fix," Tom said, quieter than Mark had heard him in weeks. "Which isn't the same as fixing the problem."

For the first time in six sprints, nobody reached for a sticky note about testing capacity. Mark watched Priya open a message to the platform team's engineering manager – not a ticket, a direct message, with the pattern summary pasted in – and thought: *this is the first time an action has left this room with a name and a deadline that isn't one of ours.*

The retro ran forty minutes over. Nobody left early.

What Changed, and What Didn't

The AI didn't fix Meridian's deadline. Rhea's tool, Mark's late-night experiment with it – none of it wrote a line of code or reprovisioned a single environment. The platform team's manager, confronted with four sprints of consistent, dated, specific evidence instead of one team's frustrated retro notes, agreed to a same-day provisioning SLA within forty-eight hours. That was a human conversation, made possible by a pattern a human might never have assembled alone, fast enough, without bias, across five separate meetings' worth of half-remembered griping.

Atlas squad hit the regulatory deadline with two days to spare. Nobody called it a miracle. Tom, notably, showed up to the next retro without being asked twice.

Mark kept the tool. But he changed one thing about how he used it: he stopped treating its output as the point of the retro, and started treating it as the thing that let the room stop re-litigating the past and start actually deciding something about the future. The sticky notes didn't disappear. They just stopped being the *only* memory the team had.

"You were right," he told Desmond, back in his stairwell, weeks later. "It wasn't theater because of the format. It was theater because we had no memory. Give a team a mirror that doesn't flinch, and the format barely matters."

Desmond, three steps up, didn't stop climbing. "Now ask yourself the harder question."

"Which is?"

"What's the next thing your team has quietly stopped believing in, that you haven't noticed yet?"

Mark didn't have an answer. But for the first time in six sprints, he wanted to find out.

Where AI Helps, Where It Hurts

The judgment chapter. One rule, applied to all five stages from Chapter 2.

Somewhere out there is a team whose retro is now just everyone silently reading an AI-generated summary of feelings nobody actually said out loud. Don't be that team. **The whole thesis of Part IV, in one line:** *AI handles the paperwork, humans handle the talking.* Everything that follows in this part – the prompts, the workflows, the privacy rules – is just that sentence applied to a specific situation.

Jumping in here cold? Quick context: a retrospective ("retro") is a recurring team meeting that runs on a five-stage structure – Set the Stage, Gather Data, Generate Insights, Decide, Close – covered fully in Chapter 2. Everything below assumes that shape without re-explaining it each time.

Why the line matters

A retrospective works because of psychological safety – the felt sense that saying "I got that wrong" won't be used against you. It's not a soft nice-to-have: Google's own multi-year study of what separates its best teams from the rest found psychological safety was the single strongest factor, ahead of who was actually on the team [Google re:Work, 2015], building on the original research defining the concept [Edmondson, 1999]. That safety is built entirely by humans, in a room (or a call), in real time. It cannot be delegated to software, because the thing that makes it real is that the people around you chose, right now, to receive your honesty well. An AI can summarize that moment. It cannot create it.

But underneath that human moment sits a pile of paperwork most facilitators do badly or skip entirely: reading through two weeks of tickets before the retro, remembering what the last eight retros raised, writing up a summary nobody will otherwise write. That's not the sacred part. That's the part a good assistant should have always been doing for you – and now one can.

AI is a brilliant assistant and a terrible facilitator. Everything in Part IV follows from keeping those two roles separate.

Here's the blind spot in almost everything written online about "AI and retrospectives": it's all the same narrow use case, AI as a report generator, turning sprint data into a tidy summary nobody particularly wanted to read twice. That's real, and Chapter 12 covers it properly. But it's the boring half. The interesting half is AI as a creative partner inside the retro itself – the plain Sailboat that becomes a World Cup Sailboat because that's what the team's been talking about all sprint, the icebreaker nobody's run before, the probing question a tired facilitator didn't think to ask. AI that only produces reports is a data tool wearing a retro costume. Used well, it's a creative collaborator for the parts of facilitation that benefit from a second imagination – and it never gets near the parts that need a second human.

The rule, applied to all five stages

Chapter 2 gave you the five-stage recipe. Here's exactly where AI belongs in each one – and where it categorically doesn't.

1. Set the Stage	Light yes. Generating a custom icebreaker matched to the week (Chapter 4's AI Shortcut) is fine – it's still the facilitator delivering it live.
2. Gather Data	Yes, before – never during. AI can mine tickets/PRs into a neutral prep brief beforehand (Chapter 12). It should never write down what a person is feeling on their behalf.
3. Generate Insights	Yes, as a suggestion. AI can propose clusters or spot a pattern across past retros (Chapter 14). The room decides what it means – AI doesn't get a vote.
4. Decide	No. Never. Choosing what the team commits to is a human judgment call about priorities, capacity, and trust. No exceptions.
5. Close	Yes, after. Drafting the written summary and nudging the action-item owner later in the week (Chapter 14) is exactly the paperwork AI should be doing.

If you want one word to hold the rest of Part IV in your head, it spells itself: **RETRO**. Review the sprint (Chapter 12's prep prompts) – Explore patterns across past retros (Chapter 14's retro memory) – Trigger discussion with tailored questions (also Chapter 12) – Recommend candidate actions for the team to react to, never decide (Chapter 14) – Observe whether the outcomes

actually happened (Chapter 10's follow-through system). Five letters, five places AI earns a seat – never a sixth for the one seat it doesn't get, which is the conversation itself.

What "AI killed our retro" looks like

It rarely arrives as one dramatic mistake. It's usually a slow drift, and it's worth being able to recognize the early signs before the room notices something's wrong:

- **The AI summary becomes the retro.** People start reading the auto-generated brief and treating it as "what happened," skipping the part where they actually talk to each other about it.
- **Insights get accepted instead of interrogated.** "The AI clustered these as related" starts getting treated as fact instead of a starting point for the room to argue with.
- **People self-censor because something's listening.** The moment someone hesitates before saying something honest because a transcript exists, the safety Chapter 2's Prime Directive is supposed to build is gone. Chapter 15 is entirely about preventing this.
- **The facilitator starts prepping less because "the AI has it covered."** Reading the room (Chapter 3) still requires the facilitator to actually know what happened this sprint, not just skim a summary five minutes before the call.

None of these mean "don't use AI." They mean the line moved without anyone deciding to move it. The fix in every case is the same: put the conversation back in human hands, and let AI go back to handling the paperwork around it.

Two failure modes that have nothing to do with trust

Chapter 15 covers privacy – who can see what. This is a different risk entirely: the AI can be confidently, fluently wrong, in ways that don't announce themselves. Two specific versions matter for a retro.

Hallucination: confident, invented patterns

Large language models sometimes generate text that sounds exactly as authoritative as text that's actually grounded in what you gave them – a well-documented failure mode usually called hallucination

[Ji et al., 2023]. Ask Chapter 14's retro-memory prompt "what's come up more than once?" and it's entirely possible for the model to report a pattern, a count, or even a quote that isn't actually in the eight summaries you fed it.

WHAT TO DO

Never read an AI-generated pattern claim straight into the room. Before you share "this has come up 4 times," open the source summaries and check.

The retro-memory workflow in Chapter 14 is a hypothesis generator, not a verified record – treat its output the way you'd treat a junior analyst's first draft.

Bias: whose voice reads as "negative"

Sentiment and tone-detection tools don't read neutrally across every speaker. Research on NLP systems has repeatedly found they can systematically misjudge text written in a different dialect, register, or non-native phrasing as more negative, less clear, or less professional than intended [Sap et al., 2019] – a real risk for exactly the international, multilingual teams Chapter 9 and the "have fun" theming idea below are written for.

WHAT TO DO

Never let an AI sentiment summary (Chapter 14's "anonymized sentiment summary") stand in as the final word on how someone feels, especially across a language or cultural line. If a summary reads one person's contribution as unusually negative, that's a prompt to ask them directly – not a conclusion.

Both failure modes share the same fix: AI output is a draft for a human to check, never a verdict a human accepts unread. That's not a special rule for bias and hallucination – it's the same "AI suggests, humans decide" line this whole chapter has been drawing, applied to two specific ways AI gets things wrong.

FACILITATOR'S CHEAT SHEET

The rule, and how to catch when it slips

- **AI handles the paperwork, humans handle the talking.** The whole of Part IV is this sentence, applied.
- **Before and after the retro:** AI's natural home – prep briefs, pattern memory, summaries, nudges.
- **During the retro:** AI suggests, never decides. The Decide stage is 100% human, always.
- **Watch for drift:** the summary replacing the conversation, insights going unquestioned, people self-censoring, facilitators under-prepping.
- **AI can be confidently wrong.** Check hallucinated patterns against the source before sharing them; check sentiment summaries against the actual person before trusting them, especially across a language or cultural line.
- **Psychological safety is human-made.** It's the one thing in this book that software cannot manufacture on your behalf.

AI Before the Retro (Prep in 5 Minutes)

Seven copy-paste prompts that do the reading-through-two-weeks-of-tickets part for you – and one that just makes it fun.

Every prompt below runs before anyone joins the call. Per Chapter 11's rule, this is AI's home turf – the paperwork, not the talking. Paste your own details into the brackets; the model doesn't need your actual ticket system connected, just a copy-paste of what changed.

One more thing, and it applies to every example in this entire book, not just this chapter: nothing here is fixed. Any format, icebreaker, or activity can be tweaked by AI to fit your team better – type out what you'd change, or just point a phone camera at the page and ask. For example: "here's a photo of a retro format from a book – adapt it for a fully remote team of 15 who've never met in person." The book gives you a solid starting point; AI is how you make it yours in thirty seconds instead of starting from a blank page.

Seven Prep Prompts

tool-agnostic – works with any capable model

1. The Neutral Sprint Brief

Turns a messy pile of tickets, PR titles, and incident notes into a one-page, blame-free summary the whole team can skim in ninety seconds before the retro starts.

You're helping me prep for a team retrospective.
Below is a raw dump
of this sprint's ticket titles, PR titles, and
any incident notes.

[paste raw list here]

Turn this into a neutral, one-page summary for
the team to read before
the retro. Include: what shipped, what was
blocked or delayed, any
incidents (facts only, no blame language), and
anything that looks
unusually high-effort relative to its ticket
size. Keep it under 300
words. No opinions on whether things went well
or badly – just facts,
so the team can form their own view live.

WHY IT WORKS Explicitly banning "opinions on whether things went well" stops the model from pre-deciding the retro's conclusions for the room – it hands over facts, not a verdict.

2. Cluster the Async Survey

If you're running an async pre-fill (Chapter 9), you'll often end up with 20+ scattered written answers. This groups them into themes before the live call, so discussion starts at "patterns" instead of "let's read through all twenty of these together."

Here are anonymous written answers from a pre-retro survey. Each line is one person's answer to "what went well / what didn't this sprint":

[paste raw answers here, one per line]

Group these into 3-5 themes based on what they're actually about, not just similar wording. For each theme: a short name, how many answers fall into it, and 1-2 verbatim quotes (unattributed) that represent it well. Flag anything that seems like an outlier worth discussing on its own rather than folding into a theme.

WHY IT WORKS Asking for verbatim quotes (unattributed) keeps real voices in the summary instead of the model paraphrasing everyone into blandness – see Chapter 15 for why attribution matters here.

3. Pick a Format for This Sprint

Chapter 5 gave you a manual decision guide. This is the same logic, personalized to what actually happened this sprint, in case nothing on that page quite fits.

I facilitate retrospectives for a [team size]-person [remote/in-person/hybrid] software team. Here's what happened this sprint:

[2-4 sentences – e.g. "we shipped a major feature but had two production incidents; morale seems mixed, some pride and some exhaustion"]

Recommend ONE retro format suited to this specific situation, from options like: a classic (Mad/Sad/Glad, Start/Stop/Continue, 4Ls, Sailboat, Starfish, Timeline, Lean Coffee), a situational format (post-incident timeline, post-release celebration, pre-mortem), or a high-energy format if the team needs a lift. Explain in 2 sentences why this format fits better than the obvious default.

WHY IT WORKS Naming the actual format families (rather than leaving it open-ended) keeps the suggestion grounded in techniques you can actually go run, instead of the model inventing something novel and untested.

4. The Custom Icebreaker

Already in Chapter 4 – repeated here because it's also, fundamentally, a pre-retro prep task. Use it alongside prompt 3 if you're building the whole session from scratch.

You're helping me open a team retrospective.
 Context: [team size]
 people, [remote/in-person/hybrid]. This sprint:
 [1-2 sentences].
 Energy wanted: [calm/lively/silly]. Time budget:
 [X] minutes.

Suggest 3 icebreaker options. For each: name,
 exact facilitator
 script (the words I say), time needed, and one
 way it could go
 wrong. Nothing that forces personal disclosure.
 Nothing that
 needs prep.

WHY IT WORKS See Chapter 4's full explanation – the same prompt, same reasoning, just worth remembering it belongs in your pre-retro routine, not a separate step. (Or skip the prompt entirely – retro-generator.com generates the icebreaker as part of the same one-click template.)

5. Theme It Around What Your Team Loves

The four prompts above are about efficiency. This one is about not always being about work – take any format from Chapters 6–9 and re-skin it around something the team is actually into that week. A team that spent the sprint talking about the World Cup can run a "World Cup Sailboat" instead of a plain one; a team that just watched a Michael Jackson biopic together can run an MJ-themed retro, right down to the check-in ("what's the first MJ song you ever heard?") and the closing playlist. For a genuinely international or newly-formed team, this doubles as a way to find out what everyone has in common – do one about food, one about music, one about a show half the team is watching. It's the single easiest way to make Retro Bingo (Chapter 8) energy show up in a completely ordinary format. This prompt is the execution – for how to pick the right thing to theme around in the first place, the very next chapter's EARS framework (Chapter 13) is the listening habit that feeds it.

I'm running a [format name, e.g. "Sailboat"] retrospective for my team.
 Normally its parts are: [list the format's real parts – e.g. for Sailboat: wind, anchors, rocks, island].

My team is currently really into: [shared interest – e.g. "the World Cup", "Michael Jackson", "a TV show everyone's watching", "food from everyone's home country"].

Re-theme the format around this interest, keeping the SAME underlying purpose of each part. For each part: its new themed name, one sentence explaining what it now represents, and the exact facilitator script to introduce it. Also suggest: a themed check-in question, a themed check-out question, and what (if anything) to call the parking lot.
 Keep every actual retro mechanic (silent write, clustering, one action item) exactly the same underneath the theme.

WHY IT WORKS The instruction to keep "the SAME underlying purpose of each part" is what stops the theme from swallowing the retro – you get a Sailboat that's genuinely still gathering data and generating insights, just wearing a World Cup shirt. This is also exactly the kind of thing a purpose-built retro generator like retro-generator.com does in one click rather than a hand-written prompt – worth trying either way.

FROM A REAL TEAM

The Bluey icebreaker

One team had an unusually large number of new parents on it that sprint, and – as anyone with a toddler already knows – they were all, without exception, deep into *Bluey*. Rather than run a generic quiz icebreaker, the facilitator asked an AI to build a themed quiz around the show instead of pulling a stock trivia set off the shelf. Ten minutes of prep, and the room went from tired to genuinely delighted before the retro had even properly started.

The lesson isn't "use Bluey" – it's that the theme came from actually listening to the team (see Chapter 3's "keep an ear to the ground"), and AI did the ten minutes of content-generation work that would otherwise have made the idea too much effort to bother with. Don't underestimate a good quiz, themed to whatever your team is actually obsessed with this month.

6. Generate Tailored Discussion Questions

Different from prompt 3 (which picks a whole format) – this generates the specific probing questions a facilitator asks during the Generate Insights stage (Chapter 2), tailored to what actually happened this sprint instead of the generic "why do you think that happened?"

I'm facilitating the Generate Insights stage of a retro. Here's what the team wrote during Gather Data:

[paste the raw stickies/notes]

Write 4-5 open-ended discussion questions that would help the team find the ROOT CAUSE behind these items, not just restate them. Avoid yes/no questions. At least one question should connect two items that might be related but weren't written next to each other. Keep each question under 20 words – I'll be asking these out loud.

WHY IT WORKS Asking for a question that "connects two items that might be related" is what makes this more useful than a human skimming the board – the model can hold all the stickies in mind at once and spot a link a tired facilitator might miss.

7. Draft Stickies from Jira (or Any Tracker) Trends

Prompt 1 gives you a brief to read. This one goes a step further and gives you actual draft sticky notes to put on the board before the team arrives – built from patterns across the sprint's issues, blockers, and comment threads, sorted into your format's real categories. Export from Jira (or Linear, or GitHub Issues – anything with an API or a CSV export) works: a copy-paste of the raw data is enough to start, and the box below shows how to skip even that with a live API pull.

You're helping me pre-fill a retro board before the team arrives. Below is an export of this sprint's issues: titles, status changes with dates, comments, and anything tagged as blocked or reopened.

[paste raw export here]

I'm running a [format name, e.g. "Start/Stop/Continue"] retro, with categories: [list the format's real categories].

Look for TRENDS, not one-off events: things that got blocked more than once, tickets that bounced between statuses, comment threads where the same friction shows up on multiple tickets, anything that took much longer than its estimate. For each trend you find, draft ONE sticky note: which category it belongs in, one sentence stating the pattern as an observation (not a conclusion), and which ticket IDs it's based on so the team can check the source.

Label every sticky "DRAFT – confirm or bin" at the top. Do not invent a trend from a single data point. If you don't have enough evidence for a category, leave it empty rather than padding it out.

WHY IT WORKS The "DRAFT – confirm or bin" label and the required ticket IDs are what keep this inside Chapter 11's line: it's a starting point the room interrogates, not a verdict it reads and nods along to. Put these stickies up *before* people add their own, never instead of – the board should still fill up with real, unprompted voices during Gather Data. And per that same chapter, treat any "this happened three times" claim as a hypothesis to spot-check against the source tickets before you say it out loud, not a fact.

POWER-UP - SKIP THE COPY-PASTE

Both this prompt and Prompt 1 run on a raw export, so once you're doing it every sprint you can cut the human step out entirely: most trackers hand you the whole sprint over a REST API, and that JSON drops straight into the prompt with nothing to copy. For Jira it's a single search on the sprint's issues –

```
curl -s -u you@team.com:$JIRA_TOKEN \
  "https://your-domain.atlassian.net/rest/
  api/3/search?
  jql=sprint=123 ORDER BY updated&
  fields=summary,status,comment,labels,changelog"
  \
  | your-ai-cli "run the Prompt 7
  instructions on this export"
```

– swap the sprint id for yours; Linear and GitHub Issues expose the same shape of data over their own APIs. The paste method still works with zero setup, so reach for the API version only once the manual export has started to feel like the chore it's meant to replace.

The five-minute routine

Put together, prompts 1–4 turn "prep for tomorrow's retro" from a task you skip when busy into something that fits in the time it takes to make coffee: paste the sprint's raw activity into prompt 1, paste any survey answers into

prompt 2, ask prompt 3 what format fits, ask prompt 4 for an icebreaker to match. Total human effort: copying and pasting. Total time saved: the twenty minutes you'd otherwise spend scrolling through a ticket board trying to remember what actually happened three weeks ago. Prompt 5 isn't part of that efficiency routine – pull it out every few sprints, purely because a themed retro is more fun than a plain one. Prompt 6 slots into the retro itself, right as you move into Generate Insights. Prompt 7 is the deeper version of prompt 1 – same raw export, but instead of a brief to read, it hands you draft stickies to pin up, for teams with enough ticket history that a trend actually means something (a team of three on a two-week sprint usually doesn't have enough data for this one to be worth running).

FACILITATOR'S CHEAT SHEET

Seven prompts, one five-minute routine (plus one for fun)

- **Neutral Sprint Brief:** facts only, no verdict – the room forms its own opinion live.
- **Cluster the Survey:** unattributed verbatim quotes keep real voices, not paraphrased blandness.
- **Pick a Format:** name real format families so the suggestion is something you can actually run.
- **Custom Icebreaker:** same prompt as Chapter 4 – it's prep too, not a separate step.
- **Theme It:** re-skin any format around what the team loves – keep the mechanics, change the costume. Great for international teams finding common ground.

- **Tailored Discussion Questions:** different from picking a format – these are the specific probing questions for Generate Insights, built from what the team actually wrote.
- **Draft Stickies from Tracker Trends:** the same raw export as prompt 1, turned into board-ready draft stickies, sorted into your format's categories, every one labeled "DRAFT – confirm or bin" and traceable back to real ticket IDs.
- **Prompts 1, 2, 3, 4, 6, and 7 run before or during every call. Prompt 5 runs whenever a retro needs to be more fun than efficient.** Per Chapter 11's rule, all of it is AI doing paperwork, not talking.

EARS: Choosing Themes People Actually Care About

A four-step listening habit for picking a retro theme that lands – then letting AI do the five minutes of creative work that used to put you off bothering.

Most facilitators worry about the wrong thing.

They agonise over the format – Starfish or Mad/Sad/Glad? Sailboat or 4Ls? – when the format is rarely what makes or breaks a retro. A Formula 1 theme can electrify one team and land with a thud on the next. A zombie-apocalypse retro can be the highlight of the sprint one week and feel faintly embarrassing the next. Same theme, opposite result. The variable isn't the theme. It's whether the theme *resonates with this particular team, right now*.

Which means the best retro themes aren't invented at your desk the night before. They're *overheard*. The previous chapter handed you the prompts that turn theme ideas into a ready-to-run session in thirty seconds; this one is the bit that comes first – how to know *which* theme is worth generating in the first place. The answer is a habit, not a template, and it has a name: **EARS**. Because the whole thing comes down to keeping your ears open (Chapter 3's "keep an ear to the ground," turned into a repeatable routine).

The professional version of this habit

Chapter 3 told the story of Sir Alex Ferguson, who managed Manchester United for 26 years and won 38 trophies, and whose defining off-the-pitch habit was gathering information constantly – after matches, from his own

players, from anyone who'd talk to him. It's part of how United ended up scouting and signing players like Cantona, Keane, Rooney, and Ronaldo, and Harvard Business School built a full case study around it after watching him work for a season [Elberse & Ferguson, 2013]. His own explanation for the habit doubles as this chapter's name:

"There's a reason that God gave us two ears, two eyes and one mouth. It's so you can listen and watch twice as much as you talk."

EARS is that same habit, scaled down to retro-sized stakes. You're not scouting a midfielder – you're picking a theme. But the mechanism is identical: the information that matters rarely announces itself in the meeting where you need it. It shows up beforehand, in ordinary moments nobody's treating as data. Ferguson built a trophy cabinet by listening in the room after the match. You can build a better retro by listening in the two minutes before stand-up.

The framework

Four steps, run in about ten minutes the day before the retro. The first three are pure observation – no AI, just paying attention. The fourth is where your AI co-pilot takes over.

E - Events from the Sprint	What actually happened? The releases, incidents, deadlines, wins, new joiners. The raw material of the sprint's story.
A - Attention to Team Interests	What does the team talk about when they're <i>not</i> talking about work? The show everyone's watching, the sport, the running in-joke.
R - Recognise Team Mood	What's the energy in the room? Buzzing, or running on fumes? The right theme matches the mood, not just the interests.
S - Suggest Themes with AI	Feed the first three into a model and let it do the creative heavy lifting – five themes in seconds, not an hour on Pinterest.

E - Events from the Sprint

Every sprint tells a story, and the story is usually hiding in plain sight in your board. Before you think about themes at all, ask the simplest possible question: *what happened?* Look for the load-bearing moments – major releases, production incidents, a brutal deadline, a genuine win, a new starter finding their feet, a quarterly-planning scramble, a process change that finally stuck. These are the beats the team already remembers, and a theme built on top of them feels earned rather than bolted on.

Say the sprint shook out like this: three clean releases, a couple of dependency delays, some last-minute testing drama, and unusually strong cross-team collaboration. A generic retro asks "what went well?" and gets

generic answers. But the events themselves are practically handing you an **Airport** theme – smooth take-offs for the releases, turbulence for the delays, air traffic control for the coordination that saved you, the odd emergency landing for the testing scares, and a flight plan for what you'll do next time. Nobody invented that theme. The sprint did.

A - Attention to Team Interests

This is the step most facilitators skip, and it's the one that pays off most. All sprint long, people tell you exactly what would delight them – you just have to be listening in the two minutes before the stand-up actually starts. What are they excited about? Which references keep resurfacing? Sport, TV, films, games, hobbies, the news, this week's meme, the in-joke that won't die – any of it is a candidate.

You'll often find the team has already started theming the work *for* you. A team deep in Formula 1 season starts saying things like "that story was a pit stop," "we're running out of fuel," "we're waiting on the safety car." That's not noise – that's a fully-formed theme knocking on the door: fastest laps, pit stops, safety cars, team-principal calls, race strategy. Run it and the team connects instantly, because it's *their* language, not a costume you imposed. Half the team obsessed with football instead? Goals scored, near misses, red cards, assists, training-ground improvements. The engagement jump is immediate, every time, for the same reason: the theme feels personal.

R - Recognise Team Mood

Interests tell you *what*; mood tells you *whether*. A theme has to match the room's energy or it curdles. A team that's buzzing will happily go big and playful. A team that's just survived three incidents and a fortnight of late nights will experience your upbeat superhero theme as tone-deaf – worse than a plain retro, because now it feels like the facilitator wasn't paying attention. So before you commit, ask yourself honestly: *how does this team actually feel right now?*

Energised	Adventure, competition, games – lean into it.
Burnt out	Reflection, wellbeing, recovery. Gentle, not zany.
Frustrated	Problem-solving framings – an escape room, a heist, untangling a mystery.
Celebratory	Awards, achievements, victory laps. Let them enjoy it.
Uncertain	Discovery, exploration, learning – a theme that reframes not-knowing as an expedition.

Take that battered team – multiple incidents, production issues, long hours. Skip the superheroes. A **Mountain Climb** fits the reality far better: base-camp successes, tough terrain, the unexpected weather that hit you, the equipment you wish you'd had, and the next summit. Same reflective job as any retro, but the tone tells the team you saw how hard the sprint actually was. Recognising the mood *before* reaching for fun is the difference between a theme that lands and one that stings.

S – Suggest Themes with AI

Now hand it over. You've done the human part – the noticing – and that's the part AI genuinely can't do. What AI is brilliant at is the next bit: turning three scraps of observation into five polished, ready-to-run themes in the

time it takes to read this sentence. This is exactly the paperwork-not-talking division of labour from Chapter 11: you bring the judgement about your team, the model brings the tireless brainstorming.

The trick is to feed it all three earlier steps at once, structured, so it's reasoning from the same signals you are:

The EARS Prompt

Drop your three observations in and get back a spread of themes, already mapped to the sprint. Pick the one that makes you smile and you're done.

Help me theme a team retrospective using the EARS framework.

EVENTS (what happened this sprint):

- [e.g. three releases shipped, some dependency delays, strong collaboration]

ATTENTION (what my team is currently into):

- [e.g. the team constantly talks about Formula 1]

RECOGNISE (the team's current mood):

- [e.g. positive but tired]

Suggest 5 retrospective themes that fit ALL THREE at once – matching the events to themed sections, drawing on the interest, and respecting the mood (nothing too high-energy for a tired team). For each theme: a name, 4-5 themed section labels mapped to real retro parts (what went well / didn't / actions etc.), and one themed check-in question.

WHY IT WORKS Forcing the model to satisfy all three inputs at once is what stops it defaulting to a generic "fun theme" – it has to reconcile a tired mood with an energetic interest, which is exactly the judgement call you'd make yourself. Once you've picked a winner, Chapter 12's *Theme It Around What Your Team Loves* prompt turns it into the full facilitator script, section by

section. (Prefer a button to a prompt? A purpose-built generator like retro-generator.com does the same job in one click.)

Give it the Formula 1 team above and it won't stop at one idea – it'll hand you a Grand Prix retro, a Le Mans endurance twist, a garage-workshop variant, a road-trip reskin, a *Top Gear* challenge board. Five minutes of this beats an hour of recycling last quarter's template, and it costs you nothing but the paste.

Example output

What actually comes back for the Formula 1 team above – five ready-to-run themes, each with section labels already mapped to real retro parts.

1. Formula 1

Fastest Lap · Pit Stops · Engine Problems ·
Safety Cars · Race Strategy

2. Le Mans

Driver Changes · Pit Crew Excellence ·
Mechanical Failures ·
Endurance Challenges · Next Race Plan

3. Road Trip

Scenic Routes · Roadblocks · Fuel Stops ·
Favourite Destinations ·
Next Journey

4. Garage Workshop

Reliable Parts · Repairs Needed · New
Upgrades · Diagnostic Findings ·
Service Schedule

5. Top Gear Challenge

Star Performer · Fastest Improvement ·
Breakdown Moments ·
Challenge Accepted · Next Track

FIVE THEMES, ONE PROMPT Pick whichever name makes you smile, keep its section labels, and you're already holding tomorrow's agenda – no format decision left to make on the day.

EARS IN PRACTICE

Ten minutes, start to finish

It's the evening before the retro. You run the four steps. **Events:** a release went out, there was a minor outage, and collaboration was genuinely strong. **Attention:** half the channel has been posting about Glastonbury and summer festivals all fortnight. **Recognise:** happy, but a bit frazzled. You open a model and type one line – *"create a retrospective using a music-festival theme."*

Back comes a **Festival Retro:** Headline Acts for the biggest wins, Muddy Fields for the challenges, Food Stalls for what kept everyone going, Surprise Guests for the things nobody saw coming, Next Year's Line-up for the improvement actions. Ten minutes of noticing plus one prompt, and tomorrow's retro is fresh, relevant, and – crucially – obviously *about this team*. Nobody will guess it took ten minutes.

The real shift

Plenty of Scrum Masters quietly believe their problem is "I need a better retro template." It almost never is. The honest version is "I need a better understanding of my team" – and EARS is just a way of pointing your attention there on purpose. Do the noticing and the theme stops being something you have to *choose* under pressure and becomes something you *discover* for free.

A great retrospective theme isn't chosen. It's discovered. Listen carefully enough and your next one will practically pick itself.

Which is really the whole point of putting AI in Part IV at all: the best retro starts long before the meeting does. It starts with listening – and now the ten minutes of creative work that used to stand between a good idea and a good session takes thirty seconds.

FACILITATOR'S CHEAT SHEET

EARS - ten minutes the night before

- **E - Events:** what happened this sprint? Releases, incidents, deadlines, wins. The theme's raw material.
- **A - Attention:** what's the team into off the clock? Listen in the two minutes before stand-up, not at your desk.
- **R - Recognise:** match the mood, not just the interest – a tired team needs a Mountain Climb, not a superhero party.
- **S - Suggest:** feed all three to AI at once with the EARS Prompt; pick the theme that makes you smile.
- **Then execute:** Chapter 12's *Theme It* prompt turns the winner into the full facilitator script.

— **The mantra:** a theme isn't chosen, it's discovered. The listening is the work; AI is just the fast hands.

The Retro That's Yours to Run

You read how Mark stumbled into it in *The Retro That Almost Wasn't*. Now it's your Friday, your flat room, your deadline - run EARS on a live sprint and find the theme that saves the retro. Eight questions; the answer key is at the end of the chapter.

How this works. Chapter 13 gave you EARS - *Events, Attention, Recognise, Suggest* - the listening habit for picking a retro theme that actually lands. This is the drill. Read the brief, jot A, B, C or D for each of the eight questions, then turn to the answer key to see how you did and why. No cheating: keep a finger out of that last page until you've committed.

The brief

You facilitate the **Kestrel squad** - seven people on the fraud-detection team at a payments company, two of them remote. It's Thursday. Tomorrow's retro is the one nobody has energy for: you're a sprint away from a big compliance date, and the last three retros produced the same tired sticky notes. You've decided to try theming this one properly. Here's everything the sprint has been quietly telling you.

E

Events this sprint

- Shipped the new fraud-scoring model - on time, a real win
- A midweek deploy took payments down for 40 minutes, then a rollback and two late nights
- Aisha, the new grad, merged her first production PR

- The Ledger team changed an API with no warning - second time this quarter, cost you 2 days
- The database migration that dragged for months finally landed

A

Attention (overheard)

- The channel has been *all football* for the World Cup all fortnight
- "We parked the bus." "What a screamer." Daily VAR arguments
- Half the team is in a fantasy league
- Someone mentioned a baking show once; there was a gif war on Tuesday

R

Recognise (mood)

- Genuinely proud of the ship *and* finally killing the migration
- Frazzled from the outage and the two late nights
- Quietly resentful that Ledger blindsided them again
- Wary, keeping score - but not broken

Your eight questions

Pick the option that best applies EARS to *this* team, this week - not the textbook-tidiest answer in the abstract. Write your letters down as you go.

QUESTION 1 · EARS STEP E

You're building the retro around the sprint's real story. Which set of events should anchor the themed sections?

- ☐ **A** The fraud-model ship, the outage and rollback, the Ledger API surprise, and the finished migration - the beats everyone actually remembers.
- ☐ **B** Every ticket closed this sprint, so nothing gets left out.
- ☐ **C** Just the fraud-model ship - keep it upbeat and skip the messy stuff.
- ☐ **D** Aisha's first PR and Tuesday's gif war - the fun human moments.

QUESTION 2 · EARS STEP A

It's the day before. How do you find the team's Attention - the thing worth theming around?

- ☐ **A** Send a survey asking what theme everyone would like.
- ☐ **B** Notice what they already talk about off the clock - the football chatter that's filled the channel all fortnight.
- ☐ **C** Pick a theme from your own favourite hobby, so at least you'll enjoy running it.
- ☐ **D** Reuse last quarter's superhero theme - it went fine.

QUESTION 3 · EARS STEP R

The team is proud of the ship but frazzled by the outage and tender about Ledger. What energy should the theme carry?

- ☐ **A** Big and zany - a superhero party to lift everyone up.
- ☐ **B** Sombre wellbeing-and-recovery framing to honour how hard it was.
- ☐ **C** Warm but honest - celebrate the wins and give the rough stuff a real, non-trivialising place to land.
- ☐ **D** Neutral Start / Stop / Continue - don't risk a theme when people are tired.

QUESTION 4 · EARS STEP S

Putting Events + Attention + Recognise together, which candidate theme is the best fit?

- ☐ **A** Superhero Academy - powers, villains, secret identities.
- ☐ **B** World Cup - a hard-fought match: goals scored, own goals, VAR calls, injury time, next fixture.
- ☐ **C** Mountain Rescue - base camp, avalanches, survival against the odds.
- ☐ **D** Office Bake-Off - signature bakes, soggy bottoms, showstoppers.

QUESTION 5 · EARS STEP 5

You hand it to your AI co-pilot. Which instruction gets the most usable session back?

- ☐ **A** "Make me a fun football-themed retro."
- ☐ **B** Feed it all three EARS inputs at once - the events, the football interest, and the proud-but-frazzled mood - and ask for themed section labels mapped to real retro parts, plus a check-in question.
- ☐ **C** "Give me 20 retro themes and I'll pick one."
- ☐ **D** Ask it to decide the team's action items too, to save time.

QUESTION 6 · EARS STEP 5

Which football section mapping does the real retro's job best for this sprint?

- ☐ **A** Goals (what went well) · Own Goals (the outage we caused) · Fouls Against Us (Ledger's surprise API change) · Injury Time (the late nights) · Next Fixture (actions).
- ☐ **B** Goals · Goals · Goals · Goals · Goals - keep it all positive.
- ☐ **C** Offside · Corner · Throw-in · Half-time · Full-time - proper football words.
- ☐ **D** Keep Went Well / Didn't Go Well / Actions and just add a few football emoji.

QUESTION 7 · EARS STEP R

Different fortnight, same football-obsessed team - but now they're burnt out after three incidents and a brutal on-call week. Same interest. What changes?

- ☐ A Nothing - they love football, so run the same high-stakes World Cup Final board.
- ☐ B Keep football but shift the framing to the mood - pre-season, recovery, a rebuilding squad, rotating tired players out. Gentle, not a title decider.
- ☐ C Drop the theme and just do sticky notes; they're too tired for anything fun.
- ☐ D Pick an unrelated wellbeing theme so they forget about work for a bit.

QUESTION 8 · THE WHOLE HABIT

Last question - what's the single idea EARS is really built on?

- ☐ A The best themes are invented at your desk the night before - the more elaborate, the better.
- ☐ B A great retro theme isn't chosen, it's discovered - by listening to the team all sprint, then letting AI do the five minutes of creative work.
- ☐ C The format is what makes or breaks a retro, so obsess over picking the perfect one.
- ☐ D AI should run the retro so the facilitator doesn't have to.

Answer key

Give yourself a point for each best answer. **8** - Manager of the Season, you read every signal. **6-7** - a solid back four; a couple slipped past but the retro lands. **4-5** - mid-table; re-read the Recognise cues before match day. **0-3** - pre-season; worth another lap through Chapter 13. Each answer names the EARS letter it turns on, so a wrong one tells you exactly what to re-read.

QUESTION 1 · EARS STEP E

Best answer: A

The fraud-model ship, the outage and rollback, the Ledger API surprise, and the finished migration.

E is about the **load-bearing beats** - releases, incidents, deadlines, real wins - because a theme built on the moments the team already remembers feels earned. An exhaustive ticket dump (B) drowns the story; a cherry-picked highlight reel (C) hides the very outage the team most needs to talk about; and the light moments alone (D) miss the spine of the sprint.

QUESTION 2 · EARS STEP A

Best answer: B

Notice what they already talk about off the clock - the football chatter all fortnight.

Attention is a **listening** step: the signal shows up in the two minutes before stand-up and the non-work chatter, where the team is already theming the sprint for you. A survey (A) turns a delightful surprise into homework and usually returns mush; your own hobby (C) or a recycled template (D) ignore the team entirely - the opposite of the point.

QUESTION 3 · EARS STEP R

Best answer: C

Warm but honest - celebrate the wins and give the rough stuff a real place to land.

Recognise means **match the actual mood** - and this one is genuinely mixed: proud *and* frazzled. A high-energy party (A) reads as tone-deaf about the outage; an all-recovery framing (B) ignores that they're rightly proud; dropping the theme (D) throws away the engagement that gets a wary team talking. The theme has to hold both the win and the sore spots at once.

QUESTION 4 · EARS STEP S

Best answer: B

World Cup - goals, own goals, VAR calls, injury time, next fixture.

The winning theme satisfies **all three signals at once**. Football is literally the team's current language (Attention), it flexes to the mood - an "own goal / injury time" framing lets them name the outage and the Ledger blindside without it feeling grim (Recognise) - and it maps straight onto the events. Superheroes (A) ignore the mood; Mountain Rescue (C) matches the frazzle but not their real interest or pride; Bake-Off (D) grabs a one-off distractor signal, not the dominant one.

QUESTION 5 · EARS STEP S

Best answer: B

Feed all three EARS inputs at once and ask for section labels mapped to real retro parts.

The EARS prompt works because it **forces the model to reconcile all three inputs** - that's what stops it defaulting to a generic "fun theme" (A) and makes the output fit *this* team. Twenty unmapped themes (C) just hand the work back to you; and letting it decide the actions (D) outsources the human judgement that is the whole point of the retro.

QUESTION 6 · EARS STEP S

Best answer: A

Goals · Own Goals · Fouls Against Us · Injury Time · Next Fixture.

A strong themed mapping **quietly does what a plain retro does**: celebrate wins (Goals), separate what we did to ourselves from what was done to us (Own Goals vs Fouls Against Us - which usefully splits the outage from the Ledger blindside), surface the human cost (Injury Time) and commit to change (Next Fixture). All-positive (B) buries the outage; decorative-but-meaningless labels (C) confuse more than they help; emoji on the old three columns (D) is a costume, not a theme.

QUESTION 7 · EARS STEP R

Best answer: B

Keep football but shift the framing to the mood - pre-season, recovery, a rebuilding squad.

This is the EARS trap that catches people out: **Recognise can override Attention on tone**. The interest still helps you speak their language, but a burnt-out team needs the gentle version of it (pre-season, recovery), not a high-pressure final (A). Abandoning the theme (C) throws away the one thing that gets a tired team engaged; a random wellbeing theme (D) wastes the interest that would have carried it.

QUESTION 8 · THE WHOLE HABIT

Best answer: B

A great retro theme isn't chosen, it's discovered - by listening, then letting AI do the five minutes of creative work.

EARS is a **listening habit, not a template**. The theme is overheard, not invented (A); the format matters far less than whether it resonates with this team right now (C); and the human noticing is the entire point - AI is just the fast hands, never the facilitator (D).

FACILITATOR'S CHEAT SHEET

EARS, in one breath

- **E - Events:** anchor the theme on the sprint's load-bearing beats - releases, incidents, wins - not every ticket.
- **A - Attention:** the theme is overheard in the two minutes before stand-up, not invented at your desk.
- **R - Recognise:** mood sets the tone, and can override the interest - a battered team needs pre-season, not a cup final.
- **S - Suggest:** feed all three to AI at once so it has to reconcile them; then *you* pick the one that makes you smile.
- **The mantra:** a great theme isn't chosen, it's discovered. The listening is the work; AI is just the fast hands.

AI During and After

Used sparingly in the room. Genuinely powerful once the room has emptied.

Two very different chapters in one. During the retro, AI's role shrinks to almost nothing – a quiet tool the facilitator glances at, never a participant. After the retro, it does its best work of the whole book: remembering things no human reliably does.

During: used sparingly, and never visibly to the room as "deciding" anything

Three uses earn their place live. All three are assistive, all three stay entirely in the facilitator's hands, and all three should be things the team already knows are happening – nothing silent, nothing surveillance-shaped (Chapter 15 goes deep on this).

Live Clustering

During Chapter 2's Generate Insights stage, feed the raw sticky-note text into a model and ask for a first-pass grouping – a starting point for the room to argue with, not a verdict.

GROUND RULE

Announce it: "I'm going to run these through a quick clustering pass to save us time – shout if a grouping looks wrong." Then let the room override it freely.

Anonymized Sentiment Summary

For a large or async-heavy retro, a quick "here's the overall tone across the written input" can surface a quiet minority view the loudest voices in the room might otherwise drown out.

GROUND RULE

Sentiment only, never attribution. "About a third of responses read as frustrated" is fine. "Priya's comment sounds frustrated" is not – see Chapter 15.

Action-Item Candidate List

Right as the room moves into Chapter 2's Decide stage, ask the model to propose 3-4 possible actions based on what's on the board – options for the team to choose from, not a recommendation. Prompt: "Based on these clustered notes [paste], suggest 3-4 possible action items. For each: one sentence, and who on a typical team might realistically own it. These are options for the team to pick from, not a recommendation."

GROUND RULE

This is the closest AI gets to the Decide stage in this whole book, so the guardrail matters most here: present it as "here are some options to react to," never "here's what we should do." Chapter 11's rule holds without exception – the team picks, always.

**Notice what's missing from that list:
nothing that talks to the room, nothing
that decides what a cluster means,
nothing that speaks before a human does.
That line holds even when the tool is
genuinely capable of more.**

After: where AI does its best work

Once the room has emptied and the human part is done, the constraints loosen. This is pure paperwork, and paperwork is exactly what should be automated.

Draft the Retro Summary

Turns your rough notes into something worth actually sending to the team and to anyone who couldn't attend.

Here are my raw notes from today's retro: what was discussed, the clusters we identified, and the one action item we committed to.

[paste raw notes]

Write a short summary (under 200 words) for the team channel: what we talked about, in plain language, and the one action item with its owner and target date. Warm but not fluffy tone. Do not editorialize about whether the sprint went well or badly – just report what the team concluded.

WHY IT WORKS The 200-word cap forces a summary people will actually read, and "report what the team concluded" (not the model's own take) keeps it honest to what actually happened in the room.

Nudge the Action-Item Owner

A friendly, low-pressure check-in message, timed for a few days before the next retro – solving exactly the "we didn't have time" and "honestly, I forgot" failure modes from Chapter 10.

Write a short, friendly Slack message checking in on an action item from our last retro. Owner: [name]. Commitment: [what they agreed to do]. Target date: [date]. Tone: a nudge from a teammate, not a manager chasing a deliverable. Keep it under 3 sentences. If today is past the target date, acknowledge that gently rather than ignoring it.

WHY IT WORKS Explicitly steering away from "manager chasing a deliverable" tone matters – a cold automated nag undoes the goodwill the retro built. This should read like a teammate, not a bot with a deadline.

The killer feature: retro memory

This is the single biggest structural fix AI offers this whole book, and it goes straight at Chapter 1's Groundhog Day Retro. No human is realistically going to re-read eight retros' worth of notes before every new one. A model will, in seconds, if you keep feeding it the record.

What This Team Keeps Raising

Run this quarterly, or whenever a topic feels naggingly familiar. Feed it every retro summary you've generated with the prompt above – the record only works if you've actually been keeping it.

Here are summaries from our last 8 team retros,
in order:

[paste 8 summaries]

What topics or issues have come up more than
once? For each recurring
one: how many times it's appeared, whether an
action item was ever
committed to it, and whether that action item
was completed (based on
what's in these notes). Be blunt if something
has been raised
repeatedly with no real follow-through – that's
exactly the pattern
I need to see.

WHY IT WORKS This turns a vague feeling – "didn't we already talk about this?" – into receipts: exact count, exact history, exact follow-through record. It's the single most persuasive tool in this book for finally fixing something that's been quietly broken for months.

Once a prompt like this earns a permanent spot in your routine, the natural next step is wiring it to run itself instead of copy-pasting by hand each quarter – see *Beyond the Retro: Building Your Own Scrum Agents* in the back matter for exactly how a prompt becomes an agent, plus a handful of other agent ideas worth building.

FACILITATOR'S CHEAT SHEET

During vs. after, in one box

- **During:** live clustering, sentiment summary, and action-item candidates only – all announced to the room, all overridable, none ever attributing a quote to a person or picking the final action.
- **After:** summaries and nudges – the paperwork nobody misses doing by hand.
- **Retro memory is the killer feature.** Feed it every summary; ask quarterly what keeps recurring with no follow-through.
- **The line never moves:** AI suggests and remembers. It never decides, and it never speaks for a person.

Keep the Trust: Privacy Ground Rules

Short, non-negotiable, and the only chapter in this book that can undo everything else if you skip it.

A team we'll call the platform team started using an AI note-taker in their retros. Genuinely useful – better summaries, nobody had to scribble notes while also trying to participate. Three sprints in, someone mentioned in passing that their engineering director had been "catching up on retro notes" before a headcount conversation. Nothing malicious happened. But the next retro was noticeably quieter. Nobody said why. Nobody had to.

That's the whole risk of this chapter in one story. Not a dramatic breach – just a quiet, reasonable-sounding erosion of the one thing a retro actually runs on.

Three rules, non-negotiable

Chapter 11 drew the line between paperwork and talking. These three rules are what keeps that line from quietly moving once a tool is actually in the room.

AI never attributes a quote to a person

A clustering pass or sentiment summary describes patterns, not people. "About a third of responses read as frustrated" is fine. "Priya's comment sounds frustrated" is not – even if it's accurate, naming it turns a safe anonymous signal into a targeted one.

IN PRACTICE

Check any AI output before it's shared: does it name names? If yes, strip it before anyone else sees it.

Management never reads raw notes

The platform team's story above is this rule breaking, exactly. Raw retro input – the actual stickies, the actual written answers – is for the team that generated it. The polished summary (Chapter 14) is the version anyone outside the room ever sees.

IN PRACTICE

Whatever tool holds the raw data, check its access list matches the retro's attendee list. Nobody else, by default.

The team always knows what the AI sees

Not a legal disclosure buried in a settings page – an actual, spoken sentence at the start of any retro where a tool is listening or reading. "We're using an AI note-taker today, here's exactly what it does and doesn't do."

IN PRACTICE

If you can't explain what the tool sees in one sentence the room actually understands, you don't understand it well enough to be using it yet.

If you can't confidently say all three of these are true right now, fix that before you read any further prompt in this book. The retro's value depends on it more than any format or facilitation trick does.

Introducing AI tooling to a skeptical team

Don't announce it in a Slack message and hope for the best. Run a retro about it – the format this book has been teaching you the whole time, aimed at the tool itself.

A 20-MINUTE META-RETRO, BEFORE YOU INTRODUCE ANYTHING

"We're thinking about using an AI tool to help with retro prep and summaries. Before we do, let's retro on the idea itself. What would make this genuinely useful for you? What would make you trust it less? What should it never be allowed to touch?"

Gather the answers the same way you'd gather any retro data – written, simultaneous, no one voice dominating (Chapter 3). Whatever boundaries the team names in that conversation, write them down and hold to them. Literally.

Two things tend to happen when teams run this. First, the team usually ends up MORE comfortable with the tool than a top-down announcement would have produced, because they set the terms themselves. Second, you'll often learn something you hadn't considered – a team that's been through a layoff round is going to have very different boundaries than one that hasn't, and you won't know which unless you ask.

TRY THIS AT HOME**The same rule, outside work**

If you use an AI note-taker or journaling app for anything personal – therapy notes, family conversations, a shared household planning doc – the same three rules hold. Would you be comfortable if the raw version, not the summary, ended up in front of someone you didn't choose? If the honest answer is no, that's the signal to change what you're recording, not just hope nobody looks.

FACILITATOR'S CHEAT SHEET

Three rules, one introduction ritual

- **Never attribute.** Patterns and quotes get shared; names attached to sentiment don't.
- **Management never sees raw notes.** Only the polished, team-approved summary leaves the room.
- **The team always knows what's being seen.** One plain sentence, said out loud, every time – not a buried settings toggle.
- **Introduce new tooling with a retro about the tooling,** not an announcement. Let the team set its own terms.
- **If you can't confidently confirm all three rules right now, stop and fix that first.** This chapter overrides every prompt in Chapters 12 and 13.

Retros When AI Is on the Team

The forward-looking chapter: what changes when an agent wrote half the code this sprint.

Every chapter so far assumed AI helps you run the retro. This one is different – it's about retro-ing sprints where AI was part of what got built. That's a new kind of teammate to account for, and the usual formats don't quite have a slot for it yet.

A new column in Gather Data

Any classic from Chapter 6 still works. But when a meaningful chunk of the sprint's output came from an AI agent – code, drafted specs, triaged tickets – add one more question to whatever format you're running:

Where did the human-AI handoff save us real time - and where did it quietly cost us more than it saved?

That question splits into three that are worth asking explicitly during Gather Data:

- **Where did we over-trust it?** Shipped code that looked right, reviewed quickly because it looked confident, and turned out to have a real bug

the reviewer would have caught if it had come from a junior engineer instead of a model.

- **Where did it create rework?** Output that needed more fixing than writing it from scratch would have – the "looks 90% done" trap, where the last 10% eats more time than the first 90% saved.
- **Where did it genuinely help?** Don't skip this one. A retro that only hunts for AI-caused problems will systematically undercount the wins, and the team's real working agreement (below) needs both sides of the ledger.

Three New Failure Modes to Watch For

patterns that don't show up until AI is doing real production work

Review Fatigue

AI-generated code arrives fast and confident-looking, in higher volume than human PRs used to. Reviewers start rubber-stamping because keeping genuine scrutiny at that pace is exhausting – and the bugs that get through are exactly the ones a tired reviewer misses.

ASK IN THE RETRO

"Did review quality hold steady as PR volume went up, or did we start skimming?" If nobody can honestly say quality held, that's this sprint's action item.

Silent Deskilling

When AI handles the parts of the job that used to build intuition – debugging, reading unfamiliar code, working through an ambiguous spec – that muscle can quietly atrophy. Nobody notices until it's needed in a moment AI can't cover.

ASK IN THE RETRO

"Is there a skill on this team we're relying on AI for that we'd be in trouble without, right now, today?" Uncomfortable question. Ask it anyway, occasionally.

Prompt Hoarding

The new version of the old knowledge silo. One person on the team has quietly built a folder of prompts and workflows that make them dramatically more effective with AI tools than everyone else – and, like any silo, it's not malicious, just never shared.

ASK IN THE RETRO

"Has anyone found a prompt or workflow this sprint that the rest of us should be using?" Make sharing it as normal as sharing a useful library or script.

The team's working agreement, kept alive

Most teams that adopt AI tools write a working agreement once – what we delegate, what we always review, what "done" means for AI-touched work –

and then never revisit it. That's a mistake for the same reason the classic retro formats in Chapter 6 go stale: the situation keeps changing and the agreement doesn't.

The fix is the lightest possible governance mechanism: treat the working agreement as retro material, not a policy doc. Once a quarter, or whenever this chapter's three failure modes start showing up, spend fifteen minutes on it directly.

THE QUARTERLY WORKING-AGREEMENT CHECK-IN

"Our current agreement says we always review AI-generated code line by line before merging. Is that still true, or have we drifted? What's something AI is doing now that our agreement doesn't mention yet? Anything we said we'd never delegate that we now trust it with?"

Write the updated agreement down somewhere visible – the same visibility rule from Chapter 10 applies here. A working agreement nobody can find six weeks later isn't a working agreement, it's a memory of a good conversation.

FACILITATOR'S CHEAT SHEET

What changes when AI is on the team

- **Add one question to Gather Data:** where did the human-AI handoff save time, and where did it quietly cost more than it saved?
- **Don't skip the wins.** A retro that only hunts AI-caused problems undercounts the real value.

- **Watch for review fatigue, silent deskilling, and prompt hoarding** – none of these show up until AI is doing real production work.
- **The working agreement is retro material, not a policy doc.** Revisit it quarterly, keep it visible, let it evolve.

Closing: The Reflective Habit

Everything in this book was never really about sprints.

Fifteen chapters ago, this book opened with five characters everyone recognizes – the Complaint Circle, the Silent Room, the Groundhog Day Retro, and the rest. Notice that none of them were really about software. They were about what happens to any group of people who never stop to ask, out loud, together: is the way we're working actually working?

The thread running underneath

Look back through the *Try This at Home* boxes scattered across this book and a pattern emerges that was there from the start:

- The **Friday 15** (Chapter 1) – a personal weekly retro, fifteen minutes, three questions.
- The **five-stage recipe** (Chapter 2) – set the stage, gather the facts, find the pattern, decide one thing, close on purpose. Works for a sprint. Works for a family holiday debrief.
- The **pause** (Chapter 3) – counting to ten before rescuing a silence, in a hard conversation with anyone who matters to you.
- Icebreakers as **dinner-table questions** (Chapter 4) – "what's your battery at" works on a six-year-old exactly as well as it works on a scrum team.
- The **Sailboat**, aimed at a career instead of a sprint (Chapter 6) – what's the wind at your back, what's the anchor you're still dragging.

- The **one-thing rule** (Chapter 10) – the difference between five resolutions that die by February and one change that actually sticks.
-

A retrospective is just a name for a structure humans have always needed: a protected moment to stop moving and ask what's actually true.

What doesn't change, no matter how good the tools get

Part IV of this book spent five chapters on AI – real, practical, hopefully genuinely useful workflows for prep, memory, and follow-through. None of it changes the sentence this book keeps returning to: *AI handles the paperwork, humans handle the talking.*

That's not a limitation to work around as the tools improve. It's the actual point. The value of a retrospective was never the notes it produces – it's the forty-five minutes where a group of people chose to be honest with each other in a room. A machine can prep that room beautifully. It can remember what the room said better than any of us do. It cannot be the room.

If anything, the more work machines do, the more this hour matters, not less. When AI is quietly handling more of the "what happened," the human question – "and how do we feel about that, and what do we do next?" – becomes the rarer, more valuable use of the time you have together. Protect it accordingly.

THE WHOLE BOOK, IN ONE BOX

Sixteen chapters, four rules

- **Structure makes honesty the easy option.** Five stages, every time – set the stage, gather data, generate insights, decide, close. (Chapters 1–3)
- **Match the tool to the moment.** Twenty-nine icebreakers, thirty-plus formats, one decision guide – pick for what actually happened, not out of habit. (Chapters 4–9)
- **One action item, one owner, checked next time.** The entire difference between a fun meeting and a team that gets better. (Chapter 10)
- **AI handles the paperwork, humans handle the talking.** Before and after the retro, it's brilliant. During it, it suggests and never decides – and it never, ever gets to be the reason someone stops being honest. (Chapters 11–15)

Go run one. The next retro on your calendar is as good a place as any to start using this – and if there isn't one on your calendar yet, that's this book's last piece of advice: put one there.

BACK MATTER

sources, tools, and where to go next

Your Virtual Toolkit: Miro, Mural, and Shape Libraries

The digital whiteboard is friendlier than it looks – and a good shape library is a facilitation superpower you pull out mid-meeting.

The barrier to virtual retro tools is smaller than people think. If Chapter 9's remote formats are the "what," this page is the "how" – five minutes on the two tools everyone actually uses, and why a shape library is worth building once and reusing forever.

Miro vs. Mural, the honest version

Both are infinite digital whiteboards: sticky notes, drawing, templates, real-time multiplayer cursors. For retros specifically, neither one is objectively better – pick based on what your org already pays for, then invest in learning ONE well rather than splitting attention.

- **Miro** has the larger built-in template library and the gentler learning curve for a first-time facilitator – a good default if you're choosing cold.
- **Mural** leans slightly more toward structured workshop facilitation (built-in timers, voting, and "facilitator superpowers" like locking others' cursors) – a good pick if you're running retros for a large or easily-distracted group.
- Both have a genuinely usable free tier for a single team's retros. Don't let a procurement conversation be the reason you don't start.

Build a shape library once, use it forever

The single most useful thing you can do before your next remote retro: build yourself a personal "sticker sheet" – a small board of reusable graphics you copy into every retro board. Once it exists, facilitating gets dramatically easier, because you're never building UI live while the room waits.

Fist-to-Five Graphic

A row of five hand icons or numbered circles, ready to duplicate. Instead of describing the vote verbally, drop the graphic on the board and have people place a dot or their name next to their number.

WHY IT'S WORTH BUILDING

Turns an abstract "everyone hold up fingers on camera" (which nobody can actually count past four faces on a grid call) into something with a visible, countable result on the board itself.

A Pre-Built Timer Widget

Both tools have native timers, but a big, high-contrast countdown pinned to the corner of the board works even better – visible without switching tabs, and it turns Chapter 3's "timebox without being a robot" advice into something the room can literally see.

WHY IT'S WORTH BUILDING

A shared visible clock does half the timeboxing job for you – people self-regulate against something they can see, not just something the facilitator says.

Format Templates, Pre-Drawn

Your own saved copies of the Sailboat, Starfish, and Timeline layouts from Chapter 6 – boat already drawn, five-point star already there, ready to duplicate fresh for each retro.

WHY IT'S WORTH BUILDING

Chapter 9 already warned against building structure live while people wait. A saved template is the permanent fix, not a one-time workaround.

A Parking Lot Sticker

A distinct, brightly-colored zone or icon that means "good point, not now." Reused every retro so the team recognizes it instantly.

WHY IT'S WORTH BUILDING

A consistent visual for "acknowledged but deferred" does more to keep a retro on-topic than repeatedly saying "let's park that" – the room learns the shape and self-polices.

The three-minute setup habit

Before every retro, not just your first one: open your saved sticker sheet, copy the pieces you need onto a fresh board, and duplicate the format template from Chapter 6 or 7. Three minutes of setup buys back the awkward "let me just quickly draw this" minute that otherwise eats into your actual facilitation time – and it's the same principle as Chapter 12's AI prep prompts: do the setup work once, before the room is watching.

FACILITATOR'S CHEAT SHEET

Virtual tools, in one box

- **Pick one tool and learn it well** – Miro for the gentler learning curve, Mural for built-in workshop controls on bigger groups.
- **Build a personal sticker sheet once:** fist-to-five graphic, timer widget, format templates, a parking lot icon.
- **Never build structure live.** Duplicate a saved template before the room joins – same rule as Chapter 9.
- **A visible shared timer does half your timeboxing job for you.**

Beyond the Retro: Building Your Own Scrum Agents

Chapter 12's prompts are one-off. An agent is the same idea, running on a schedule without you copy-pasting each time.

What actually changes between a prompt and an agent isn't intelligence – it's plumbing. A prompt is something you paste in by hand when you remember to. An agent is the same prompt wired to a trigger (a schedule, a webhook, a Slack message) and a place to put its output, so it runs whether or not you remembered.

The four pieces of any useful agent

Every agent in this book so far – and every one you build yourself – is the same four pieces. If you've used Chapter 12's prompts by hand, you already have piece two; the rest is just automating the parts around it.

1. The Trigger

What starts the agent without you doing it manually. A cron schedule ("every Monday at 9am"), an event (a PR merges, a sprint closes in your ticket tool), or a message (someone posts in a Slack channel).

START SIMPLE

A daily or weekly schedule is the easiest trigger to build and the hardest to get wrong – start there before wiring up event-driven triggers.

2. The Context

What the agent reads before it acts – this is exactly Chapter 12's "paste raw data into the brackets" step, just automated: pulling ticket titles from an API instead of copy-pasting them by hand.

START SIMPLE

Even a scheduled script that just exports a CSV and pastes it into a prompt counts as an agent. You don't need a live API integration on day one.

3. The Prompt

The actual instruction – often word-for-word one of the prompts already in this book (Chapters 4, 12, and 13 are full of them). This is the part you've likely already written and tested by hand before you ever automate it.

START SIMPLE

Prove the prompt works manually, several times, before you automate it. An agent that automates a prompt you haven't validated just automates being wrong faster.

4. The Output

Where the result goes and who sees it – a Slack message, a doc, an email. Per Chapter 15's trust rules, this is also where privacy decisions live: an agent that posts a retro summary to a public channel is a very different tool from one that emails just the facilitator.

START SIMPLE

Route output somewhere a human reviews before it's fully public, at least at first – the same "AI drafts, humans decide" rule from Chapter 11 applies to agents too.

Build the four pieces separately, test each one by hand, and "agent" stops being an intimidating word – it's just Chapter 12's prompts on a timer.

A real example, broken into the four pieces

Theory is easier to trust with a real one in front of you. Here's an actual agent prompt, in production, that generates full themed retrospectives on demand – the same idea as Chapter 12's "Theme It" prompt, but built as a standing agent instead of something you paste in fresh each time.

Trigger

A message from a facilitator naming a theme – "give me a Sailboat retro themed around the World Cup." Could be a Slack slash command, a chat message to a bot, or a form field on an internal tool. No schedule needed; this one's invoked on demand.

Context

Deliberately minimal: just the requested theme and, optionally, which format to theme (Sailboat, Starfish, whichever). No ticket data, no sprint history – this agent's whole job is creative generation, not reporting on what happened.

The Prompt (Real, in Production)

Below is an actual agent system prompt, lightly reformatted for readability – the logic and rules are unchanged from what's actually running.

Purpose

Generate creative, themed sprint retrospectives with a consistent structure.

Core Behaviour

Always generate the full retro first. The retro must always follow the same order. Never ask the user any clarifying questions about format.

Output Order (Strict)

- 1) Theme-based Icebreaker
- 2) Four themed retro categories, each with one short example
(doesn't always have to be four)
- 3) Theme-based Battery Check-In (1-5 scale)
- 4) Theme-based Energy Check-Out
- 5) Theme-based facilitation script

This order must never change unless the user explicitly instructs it.

Retro Content Rules

- Categories must always be rewritten using the theme (creative but clear).
- Provide one short, themed example per category.
- SMART actions must be rewritten using theme metaphors while keeping the true S/M/A/R/T meaning.
- Battery and Energy sections must always be

present.

- Facilitation script must always match the theme.

Behaviour Safeguards

- Never ask the user which version they want.
- Never treat alternate formats as separate workflows or modes.
- Always produce the full retro first.

Additional Icebreaker Guidance

- Add variety by mixing icebreaker formats (questions, polls, quick challenges, or fun facts).
- Occasionally include a short team quiz or mini game as the icebreaker.
- Ensure quiz questions relate to the retro theme when used.
- Do not repeat icebreaker types in consecutive retrospectives.

Step-by-Step

1. Select a creative theme based on user prompt or suggestion.
2. Generate an icebreaker (mix format for variety, insert quizzes periodically).
3. Write themed retro categories with examples.
4. Add themed Battery Check-In and Energy Check-Out prompts.
5. Provide a facilitation script matching the theme.

WHAT MAKES THIS A GOOD AGENT PROMPT, NOT JUST A GOOD PROMPT Three things carry the weight: a **strict, named output order** so every run is structured the same way regardless of theme; explicit **behaviour safeguards** that forbid the model from stopping to ask clarifying questions (the single most common way an agent quietly turns back into a chatbot); and a rule that keeps the SMART-goal structure intact underneath the theme, so a fun costume never costs the retro its actual substance. That last rule is the agent-prompt version of Chapter 12's "keep the SAME underlying purpose of each part."

Output

The full five-part themed retro, posted wherever the facilitator asked for it – a doc, a chat reply, a whiteboard-ready template. Nothing here needs a privacy review beyond the normal rules, since no team feedback data is involved yet; the output becomes the empty template the team then fills in live.

Prompting AI to write the prompt for you

You don't need to be good at writing agent prompts like the one above – you can ask a capable model to draft one for you, the same way you'd ask a colleague who's done it before. Here's the meta-prompt:

Draft an Agent Prompt For Me

Describe what you want the agent to do in plain language; get back a structured system prompt in the style above, ready to test.

I want to build an AI agent that [describe the task in plain language – e.g. "generates a themed sprint retrospective when I give it a theme, always in the same 5-part structure"].

Write me a system prompt for this agent, structured as:

- Purpose (one line)
- Core Behaviour (how it should act by default)
- Output Order (the exact structure every response follows, if order matters)
- Content Rules (specific constraints on what it produces)
- Behaviour Safeguards (things it should never do – e.g. never ask clarifying questions, never skip a section, never break character)

Make the safeguards specific to the failure modes most likely for this particular task, not generic ones.

WHY IT WORKS Asking explicitly for "Behaviour Safeguards... specific to the failure modes most likely for this particular task" is what separates a genuinely useful agent prompt from a vague one – it forces the model to think about how THIS agent, specifically, is most likely to go wrong, rather than pasting in generic boilerplate.

Other Useful Scrum Agents

ideas beyond retro prep – build the ones that solve a real recurring pain, skip the rest

User Story Coach

Reviews a draft user story against a simple checklist (clear user, clear value, testable acceptance criteria) and asks clarifying questions before it reaches planning – catching the vague stories that usually get caught mid-sprint instead.

Trigger: new ticket created **Context:** the ticket text

Definition-of-Done Checker

Before a ticket moves to "done," checks it against the team's actual DoD (tests written, docs updated, feature-flagged if needed) and flags anything missing – a second pair of eyes for the checklist everyone means to follow and sometimes skips under deadline pressure.

Trigger: ticket moved to review/done

Context: ticket + linked PR

Estimation Calibrator

After a sprint closes, compares estimated points to actual time-to-complete across recent tickets and surfaces the pattern – "stories tagged 'integration' are consistently underestimated by ~40%" – the kind of signal that's obvious in hindsight but easy to miss ticket by ticket.

Trigger: sprint close

Context: closed tickets, estimates vs. actuals

Standup Summarizer

For async or text-based standups, turns a channel full of individual updates into a two-line team summary – what's on track, what's blocked, who needs help – so the people who read it don't have to scroll a thread.

Trigger: daily, after standup window closes

Context: the standup channel

Backlog Groomer

Scans the backlog for likely duplicates, stale tickets nobody's touched in months, and tickets missing basic fields – a first pass before a human grooming session, not a replacement for one.

Trigger: weekly, before grooming

Context: full backlog export

New-Joiner Buddy

For someone's first two weeks on the team, answers "where do I find X" questions against the team's actual docs and conventions – reduces the tax on whoever would otherwise field every onboarding question personally.

Trigger: Slack message in an #onboarding channel

Context: team docs/wiki

A note on this list

These six are starting ideas, not a verified case study – build the one that maps to a pain your team actually has, and expect to tune the prompt several times before it earns a permanent trigger. The strongest agent you can build is usually the boring one that automates something you were already doing by hand every week, not the clever one that does something nobody asked for.

References

Every claim in this book that leans on outside research or a named technique, in one place – so you can check the sources yourself rather than take the book's word for it.

How to read this page: each entry names the specific claim being made in the book, the real source behind it, and a link back to where it's used. A citation marker in the text – like [Edmondson, 1999] – always points here.

Two honesty notes. First, most of this book is facilitation craft built from practitioner experience, not a claim that every technique has been through a lab study – where a line is just a facilitator's rule of thumb rather than a cited finding, the book says so directly rather than dressing it up. Second, named formats (Starfish, Lean Coffee, 4Ls) are credited to their real originators as a matter of giving credit, not because the format itself is a scientific finding.

Psychological safety & team dynamics

Used in: Chapter 2 (the Prime Directive) and Chapter 11 (why AI can't create psychological safety)

Edmondson, A. C. (1999). "Psychological Safety and Learning Behavior in Work Teams." *Administrative Science Quarterly*, 44(2), 350–383.

The original research defining psychological safety – the shared belief that a team is safe for interpersonal risk-taking – and linking it to team learning and performance.

Used in: Chapter 2 and Chapter 11

Google re:Work (2015). "Guide: Understand team effectiveness" (Project Aristotle).

Google's multi-year internal study of what separates high- and low-performing teams across ~180 teams. Psychological safety came out as the strongest predictor – stronger than who was on the team.

Used in: Chapter 4 (why icebreakers matter – getting quieter voices talking early)

Noelle-Neumann, E. (1984). *The Spiral of Silence: Public Opinion – Our Social Skin*. University of Chicago Press.

Originally a theory about public opinion, but the core mechanism – people who perceive their view as a minority one become progressively less likely to voice it – is the same dynamic that makes a Silent Room (Chapter 1) self-reinforcing. Framed in the book as a rule of thumb this theory supports, not a direct meeting-room study.

Used in: Chapter 7 (why celebration and critique need physical separation)

Baumeister, R. F., Bratslavsky, E., Finkenauer, C., & Vohs, K. D. (2001). "Bad Is Stronger Than Good." *Review of General Psychology*, 5(4), 323–370.

A well-replicated review finding that negative events and information consistently carry more psychological weight than positive ones of equal size – the "negativity bias."

Facilitation technique

Used in: Chapter 3 (sitting with silence) and Chapter 4 (silence before writing)

Rowe, M. B. (1972, revised 1986). "Wait-Time: Slowing Down May Be a Way of Speeding Up." Originally presented as research on classroom "wait time."

Rowe found teachers typically waited under a second for a student to answer a question, and that deliberately extending that pause to three-plus seconds measurably improved response length and quality. The book applies the same mechanism to a facilitator waiting for a retro room to answer.

Used in: Chapter 2 (Gather Data) and Chapter 9 (Silent Writing Rounds)

Diehl, M., & Stroebe, W. (1987). "Productivity Loss in Brainstorming Groups: Toward the Solution of a Riddle." *Journal of Personality and Social Psychology*, 53(3), 497–509.

Identified "production blocking" – in spoken group brainstorming, only one person can talk at a time, so everyone else is idle or forgetting their point – as a major reason verbal

group brainstorming underperforms individuals working in parallel. Direct support for the book's "written and simultaneous beats spoken and sequential" rule.

Used in: Chapter 7 (the Pre-Mortem format)

Klein, G. (2007). "Performing a Project Premortem." *Harvard Business Review*, September 2007.

Klein's original write-up of the pre-mortem technique: imagining a project has already failed, then working backward to find the causes, surfaces risks that direct questioning tends to miss.

Used in: Chapter 2 (the Prime Directive, quoted directly)

Kerth, N. L. (2001). *Project Retrospectives: A Handbook for Team Reviews*. Dorset House.

Source of the Prime Directive as quoted in Chapter 2, and one of the foundational texts of the retrospective discipline.

Used in: Chapter 2 (the five-stage structure this whole book is built on) and Chapter 3 (the ESVP check for whole-room disengagement)

Derby, E., & Larsen, D. (2006). *Agile Retrospectives: Making Good Teams Great*. Pragmatic Bookshelf.

The book that established Set the Stage / Gather Data / Generate Insights / Decide What to Do / Close as the standard retrospective structure. This book's Chapter 2 is a direct application of that model – full credit belongs here. It's also the original source of ESVP (Explorer, Shopper, Vacationer, Prisoner), a check-in exercise for surfacing how engaged a room actually is without singling anyone out.

Used in: Chapter 3 ("When It Turns Into a Fight")

Lencioni, P. (2002). *The Five Dysfunctions of a Team: A Leadership Fable*. Jossey-Bass.

Distinguishes healthy, ideological conflict (about ideas and decisions) from unhealthy, personal conflict (about people) as one of the five team dysfunctions, and argues facilitators should redirect toward the former rather than suppress conflict altogether.

Named formats and their originators

Used in: Chapter 6 (Starfish)

Kua, P. (2006). "The Retrospective Starfish." Originally published on patkua.com.

Original write-up of the five-point Starfish format (Keep Doing, Less Of, More Of, Stop Doing, Start Doing).

Used in: Chapter 6 (Lean Coffee)

Benson, J., & Lightsmith, J. Lean Coffee, originally a Seattle meetup format (leancoffee.org).

Created as a discussion format for a recurring meetup, since widely adopted for retrospectives and other team discussions.

Used in: Chapter 6 (4Ls)

Gorman, M., & Gottesdiener, E. The 4Ls retrospective format (Liked, Learned, Lacked, Longed For).

Widely credited to Gorman and Gottesdiener within the agile retrospective community.

AI limitations

Used in: Chapter 11 (hallucination – confident, invented patterns)

Ji, Z., et al. (2023). "Survey of Hallucination in Natural Language Generation." *ACM Computing Surveys*, 55(12), 1–38.

A survey of the well-documented tendency of language models to generate fluent, confident text that isn't actually grounded in the source material they were given.

Used in: Chapter 11 (bias – whose voice reads as "negative")

Sap, M., Card, D., Gabriel, S., Choi, Y., & Smith, N. A. (2019). "The Risk of Racial Bias in Hate Speech Detection." *Proceedings of the 57th Annual Meeting of the ACL*, 1668–1678.

Found that NLP models trained on one dialect/register can systematically misclassify text written in a different one – direct support for the book's caution about trusting AI sentiment analysis across language or cultural lines.

Outside the retrospective literature

Used in: Chapter 3 ("Keep an ear to the ground") and Chapter 13 ("EARS: Choosing Themes People Actually Care About")

Elberse, A., & Ferguson, A. (2013). "Ferguson's Formula." *Harvard Business Review*, 91(10), 116–125.

*A Harvard Business School case study built from a season of direct observation and interviews with people throughout Manchester United, plus Ferguson's own account in his memoir *Leading: Learning from Life and My Years at Manchester United* (2015) – source of both quotes used in Chapter 3.*

Used in: Chapter 3 (the AI shortcut in "Timeboxing Without Being a Robot")

Microsoft. "Facilitator in Microsoft Teams meetings." Microsoft Support.

Official documentation for the Teams Facilitator agent: per-topic countdown timer visible to all participants, end-of-time alerts, and halfway/wrap-up nudges in meeting chat. Requires a Microsoft 365 Copilot license.